



CSC 348 – Intro to Compilers

Lecture 6

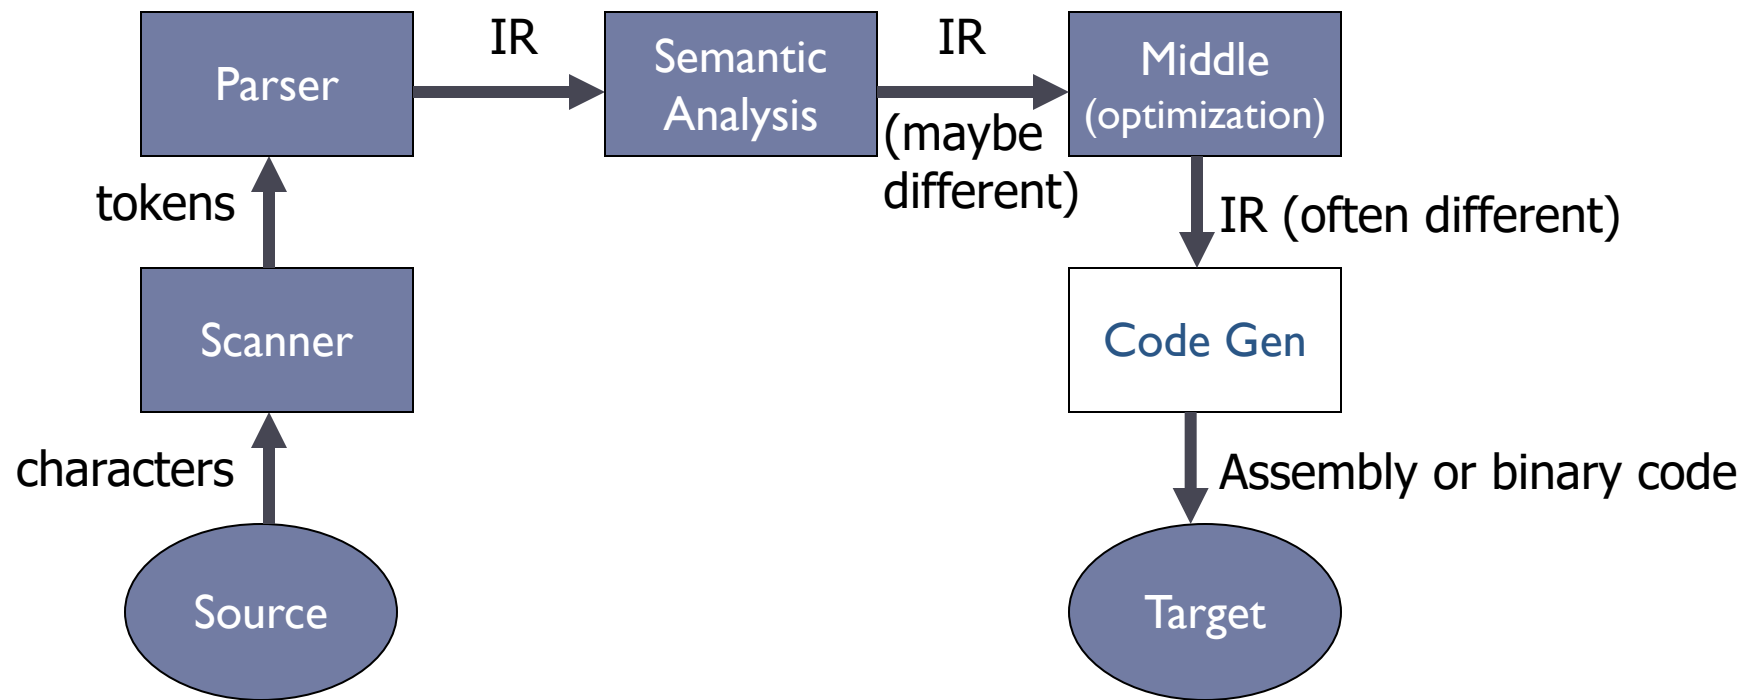
Dr. David Zaretsky
david.zaretsky@depaul.edu

Agenda



- ▶ Overview of x86-64 architecture
- ▶ Programming Assignment 4

Compiler Structure (review)



Code Translation

- ▶ **Code translation can vary based on:**
 - ▶ Different programs have different algorithms & code structures
 - ▶ Different architectures expose different interfaces
 - ▶ Hardware implementations have different performance
- ▶ **Compilers figure out what instructions to generate from the code**
 - ▶ Employ various optimizations
 - ▶ Use combination of instructions to implement high-level operations

HW Interface Affects Performance

Source code

Different applications
or algorithms

Compiler

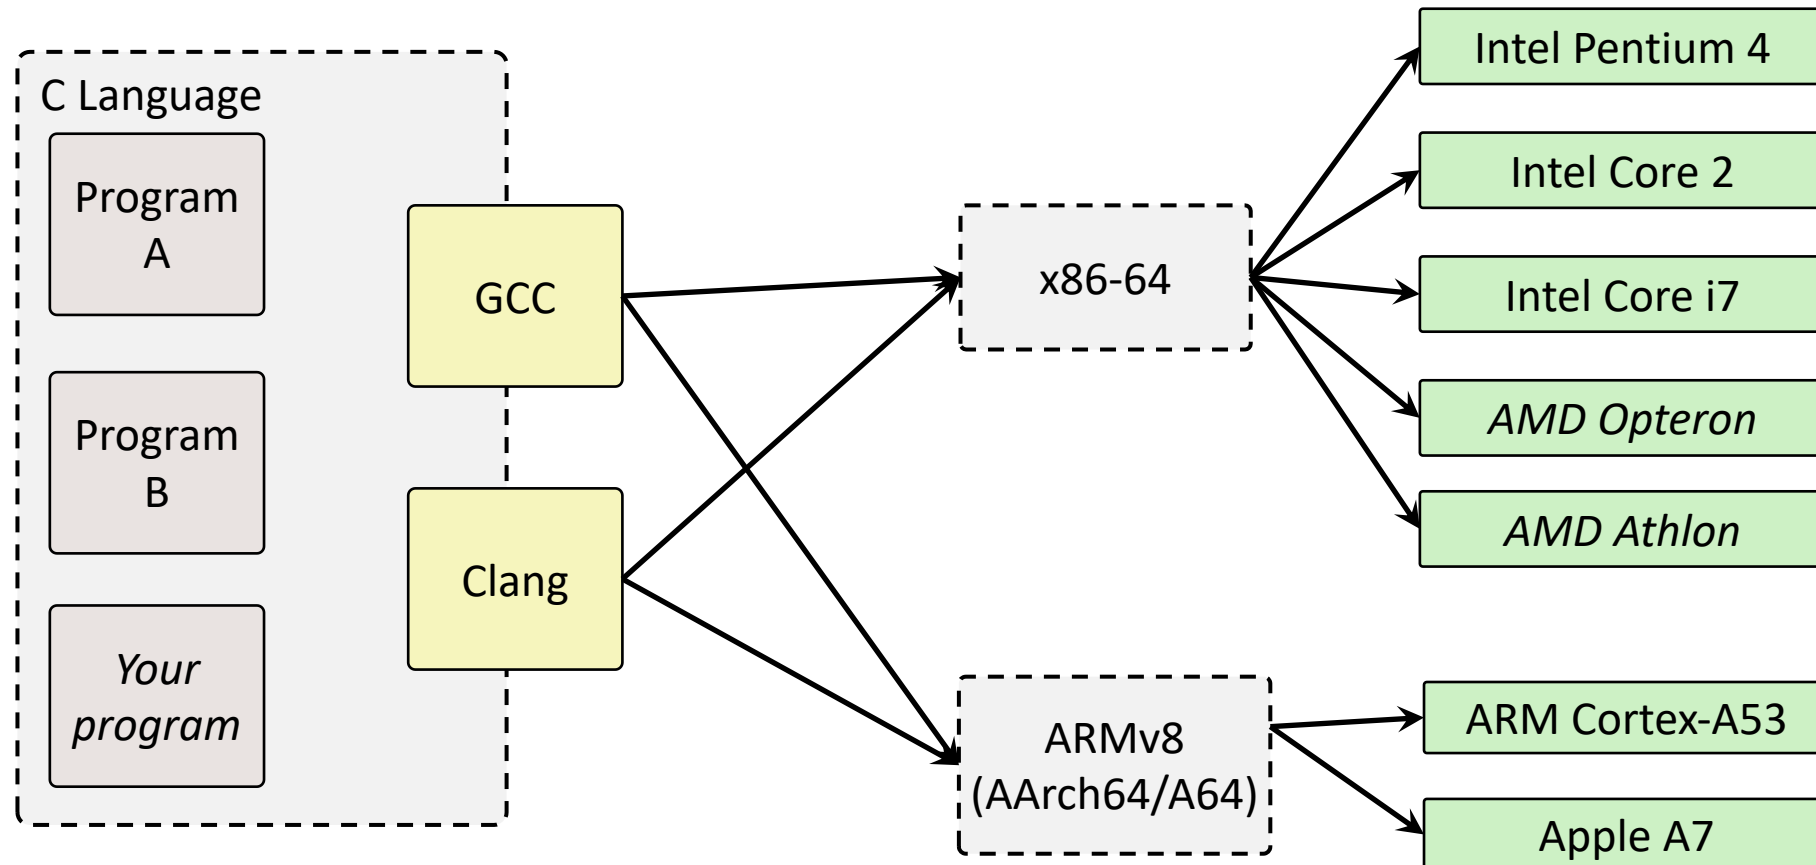
Perform optimizations,
generate instructions

Architecture

Instruction set


Hardware

Different
implementations



Instruction Set Architectures

- ▶ The ISA defines:
 - ▶ The system's state (e.g. registers, memory, program counter)
 - ▶ The instructions the CPU can execute
 - ▶ The effect that each of these instructions will have on the system state



Architecture - CISC vs. RISC

▶ Complex Instruction Set Computing (CISC):

- ▶ Add more and more elaborate and specialized instructions as needed
- ▶ Lots of tools for programmers to use, but hardware must be able to handle all instructions
- ▶ Many very special purpose instructions that you will never see, and a given compiler may never use - just need to know how to use the manual
- ▶ Variable-length instructions, between 1 and 16(?) bytes long.
 - ▶ 16 is max len in theory, I don't know if it can happen in practice
- ▶ x86-64 is CISC, but only a small subset of instructions encountered with Linux programs

▶ Reduced Instruction Set Computing (RISC):

- ▶ Typically more registers, less and fixed-size instructions
- ▶ Easier to build fast hardware
- ▶ Let software do the complicated operations by composing simpler ones
- ▶ Other major architectures are typically RISC
- ▶ Examples: PowerPC, ARM, SPARC, MIPS

General ISA Design Decisions

▶ Instructions

- ▶ What instructions are available? What do they do?
- ▶ How are they encoded?


▶ Registers

- ▶ How many registers are there?
- ▶ How wide are they?

▶ Memory

- ▶ How do you specify a memory location?
- ▶ How do you read / write to a memory location?

Mainstream ISAs



x86

Designer	Intel, AMD
Bits	16-bit, 32-bit and 64-bit
Introduced	1978 (16-bit), 1985 (32-bit), 2003 (64-bit)
Design	CISC
Type	Register-memory
Encoding	Variable (1 to 15 bytes)
Endianness	Little

Macbooks & PCs
(Core i3, i5, i7, M)

[x86-64 Instruction Set](#)



ARM architectures

Designer	ARM Holdings
Bits	32-bit, 64-bit
Introduced	1985; 31 years ago
Design	RISC
Type	Register-Register
Encoding	AArch64/A64 and AArch32/A32 use 32-bit instructions, T32 (Thumb-2) uses mixed 16- and 32-bit instructions. ARMv7 user-space compatibility ^[1]
Endianness	Bi (little as default)

Smartphone-like devices
(iPhone, iPad, Raspberry Pi)

[ARM Instruction Set](#)



MIPS

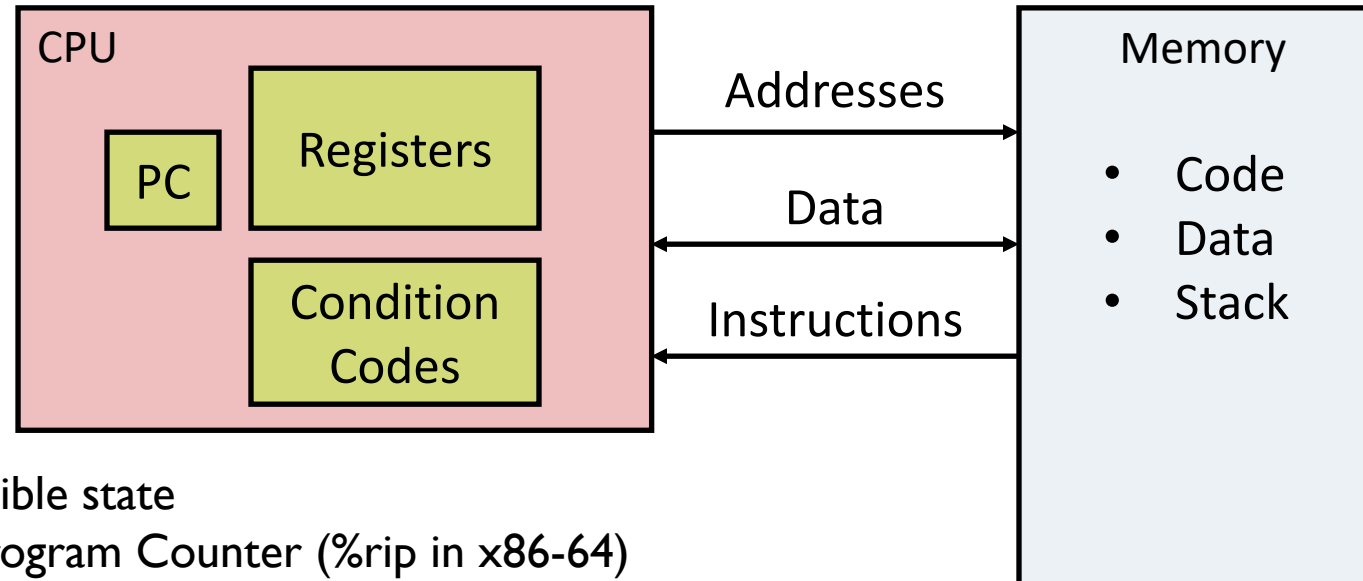
Designer	MIPS Technologies, Inc.
Bits	64-bit (32→64)
Introduced	1981; 35 years ago
Design	RISC
Type	Register-Register
Encoding	Fixed
Endianness	Bi

Digital home & networking
equipment

(Blu-ray, PlayStation 2)

[MIPS Instruction Set](#)

Assembly Programmer's View



Programmer-visible state

PC: the Program Counter (%rip in x86-64)

Address of next instruction

Named registers

Together in “register file”

Heavily used program data

Condition codes

Store status information about most recent arithmetic operation

Used for conditional branching

Memory

Byte-addressable array

Code and user data

Includes *the Stack* (for supporting procedures)

x86-64 References

- ▶ [x86-64 Instructions and ABI](#)
 - ▶ Handout for University of Chicago CMSC 22620 by John Reppy
- ▶ [x86-64 Machine Level Programming](#)
- ▶ [Quick guide to GDB](#)
- ▶ [Intel processor documentation](#)
- ▶ [GNU assembler documentation](#) ([x86 and x86-64 details](#), [AT&T vs Intel syntax](#))

x86-64 Main features

- ▶ 16 64-bit general registers
- ▶ 64-bit address space; pointers are 8 bytes
- ▶ 16 SSE registers for floating point, SIMD
- ▶ Register-based function call conventions
- ▶ Additional addressing modes (pc relative)
- ▶ 32-bit legacy mode
- ▶ Some pruning of old features

x86-64 Assembler Language

- ▶ Target for our compiler project
- ▶ But, the nice thing about standards...
- ▶ Two main assembler languages for x86-64
 - ▶ Intel/Microsoft version – what's in the Intel docs
 - ▶ AT&T/GNU assembler – what we're generating and what's in the linked references
 - ▶ Use gcc -S to generate asm code from C/C++ code
- ▶ Slides use gcc/AT&T/GNU syntax

Intel vs. GNU Assembler

► Main differences between Intel docs and gcc assembler

	Intel/Microsoft	AT&T/GNU as
Operand order: op a,b	a = a op b (dst first)	b = a op b (dst last)
Memory address	[baseregister+offset]	offset(baseregister)
Instruction mnemonics	mov, add, push, ...	movq, addq, pushq [explicit operand size added to end]
Register names	rax, rbx, rbp, rsp, ...	%rax, %rbx, %rbp, %rsp, ...
Constants	17, 42	\$17, \$42
Comments	; to end of line	# to end of line or /* ... */

- Intel docs include many complex, historical instructions and artifacts that aren't commonly used by modern compilers – and we won't use them either

x86-64 Assembly “Data Types”

- ▶ Integral data of 1, 2, 4, or 8 bytes
 - ▶ Data values
 - ▶ Addresses (untyped pointers)
- ▶ Floating point data of 4, 8, 10 or 2x8 or 4x4 or 8x2
 - ▶ Different registers for those (e.g. %xmm1, %ymm2)
 - ▶ Come from extensions to x86 (SSE, AVX, ...)
- ▶ No aggregate types such as arrays or structures
 - ▶ Just contiguously allocated bytes in memory
- ▶ Two common syntaxes
 - ▶ “AT&T”: used by our course, slides, textbook, gnu tools, ...
 - ▶ “Intel”: used by Intel documentation, Intel tools, ...
 - ▶ Must know which you’re reading

} Not covered
In CSC-348

What is a Register?

- ▶ A location in the CPU that stores a small amount of data, which can be accessed very quickly (once every clock cycle)
- ▶ Registers have names, not addresses
 - ▶ In assembly, they start with % (e.g. %rsi)
- ▶ Registers are at the heart of assembly programming
 - ▶ They are a precious commodity in all architectures, but especially x86

x86-64 Integer Registers – 64 bits wide

- ▶ Can reference low-order 4 bytes (also low-order 2 & 1 bytes)
- ▶ One is particularly important to not misuse: **%rsp – Stack Pointer**

%rax	%eax
%rbx	%ebx
%rcx	%ecx
%rdx	%edx
%rsi	%esi
%rdi	%edi
%rsp	%esp
%rbp	%ebp

%r8	%r8d
%r9	%r9d
%r10	%r10d
%r11	%r11d
%r12	%r12d
%r13	%r13d
%r14	%r14d
%r15	%r15d

x86-64 registers

- ▶ 16 64-bit general registers
 - ▶ %rax, %rbx, %rcx, %rdx, %rsi, %rdi, %rbp, %rsp, %r8-%r15
- ▶ Registers can be used as 64-bit integers or pointers, or as 32-bit ints
 - ▶ Also possible to reference low-order 16- and 8-bit chunks – we won't for most part
- ▶ To simplify our project we'll use only 64-bit data (ints, pointers, even booleans!)

X86-64 Registers

Register	Callee Save	Description
%rax		result register; also used in <code>idiv</code> and <code>imul</code> instructions.
%rbx	yes	miscellaneous register
%rcx		fourth argument register
%rdx		third argument register; also used in <code>idiv</code> and <code>imul</code> instructions.
%rsp		stack pointer
%rbp	yes	frame pointer
%rsi		second argument register
%rdi		first argument register
%r8		fifth argument register
%r9		sixth argument register
%r10		miscellaneous register
%r11		miscellaneous register
%r12-%r15	yes	miscellaneous registers

Three Basic Kinds of Instructions

- ▶ Transfer data between memory and register
 - ▶ Load data from memory into register
 - ▶ `%reg = Mem[address]`
 - ▶ Store register data into memory
 - ▶ `Mem[address] = %reg`
- ▶ Perform arithmetic operation on register or memory data
 - ▶ `c = a + b; z = x << y; i = h & g;`
- ▶ Control flow: what instruction to execute next
 - ▶ Unconditional jumps to/from procedures
 - ▶ Conditional branches

Remember: Memory is indexed just like an array of bytes!

Operand types

- ▶ **Immediate: Constant integer data**
 - ▶ Examples: \$0x400, \$-533
 - ▶ Like C literal, but prefixed with ‘\$’
 - ▶ Encoded with 1, 2, 4, or 8 bytes depending on the instruction
- ▶ **Register: 1 of 16 integer registers**
 - ▶ Examples: %rax, %r13
 - ▶ But %rsp reserved for special use
 - ▶ Others have special uses for particular instructions
- ▶ **Memory: Consecutive bytes of memory at a computed address**
 - ▶ Simplest example: (%rax)
 - ▶ Various other “address modes”

`%rax``%rcx``%rdx``%rbx``%rsi``%rdi``%rsp``%rbp``%rN`

x86-64 Memory Model

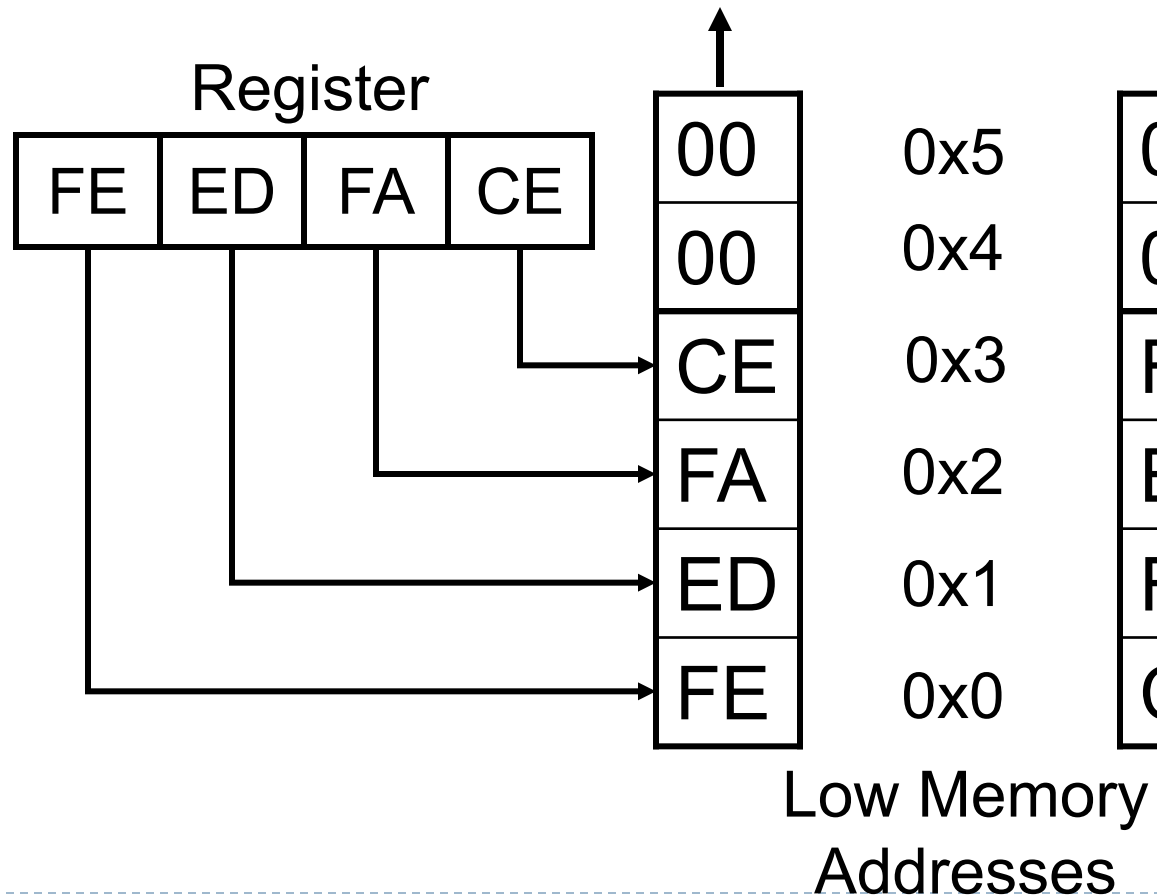
- ▶ 8-bit bytes, byte addressable
- ▶ 16-, 32-, 64-bit words, double words and quad words (Intel terminology)
 - ▶ That's why the 'q' in 64-bit instructions like `movq`, `addq`, etc.
- ▶ Data should usually be aligned on “natural” boundaries for performance, although unaligned accesses are generally supported – but with a big performance penalty on many machines
- ▶ Little-endian – address of a multi-byte integer is address of low-order byte

Architecture - Endian

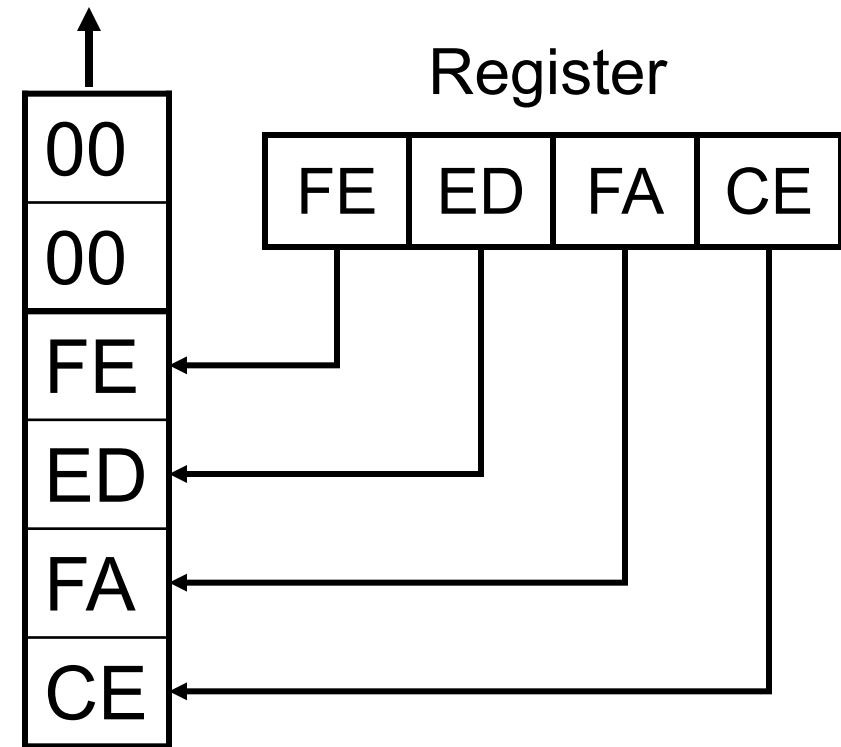
- ▶ Endianness comes from Jonathan Swift's Gulliver's Travels. It doesn't matter which way you eat your eggs :)
- ▶ Little Endian - 0x12345678 stored in RAM “little end” first. The least significant byte of a word or larger is stored in the lowest address. E.g. 0x78563412
 - ▶ Intel is Little Endian
- ▶ Big Endian - 0x12345678 stored as is.
 - ▶ Network traffic is Big Endian
 - ▶ Most everyone else you've heard of (PowerPC, ARM, SPARC, MIPS) is either Big Endian by default or can be configured as either (Bi-Endian)

Endianess pictures

Big Endian (Others)



Little Endian (Intel)



Instruction Format

- ▶ Typical data manipulation instruction

opcode src,dst # comment

- ▶ Meaning is

$\text{dst} \leftarrow \text{dst op src}$

- ▶ Normally, one operand is a register, the other is a register, memory location, or integer constant
 - ▶ Can't have both operands in memory – can't encode two memory addresses in a single instruction (e.g., cmp, mov)
- ▶ Language is free-form, comments and labels may appear on lines by themselves (and can have multiple labels per line of code)

Moving Data

- ▶ General form: `mov_ source, destination`
 - ▶ Missing letter (`_`) specifies size of operands
 - ▶ Note that due to backwards-compatible support for 8086 programs (16-bit machines!), “word” means 16 bits = 2 bytes in x86 instruction names
 - ▶ Lots of these in typical code

- ▶ `movb src, dst`
 - ▶ Move 1-byte “byte”
- ▶ `movw src, dst`
 - ▶ Move 2-byte “word”
- ❖ `movl src, dst`
 - Move 4-byte “long word”
- ❖ `movq src, dst`
 - Move 8-byte “quad word”

movq Operand Combinations

	Source	Dest	Src, Dest	C
movq	Imm	Reg	movq \$0x4, %rax	var_a = 0x4;
		Mem	movq \$-147, (%rax)	*p_a = -147;
	Reg	Reg	movq %rax, %rdx	var_d = var_a;
		Mem	movq %rax, (%rdx)	*p_d = var_a;
	Mem	Reg	movq (%rax), %rdx	var_d = *p_a;

- ❖ *Cannot do memory-memory transfer with a single instruction*
 - How would you do it?

Basic Data Movement and Arithmetic Instructions



`movq src,dst`

$\text{dst} \leftarrow \text{src}$

`addq src,dst`

$\text{dst} \leftarrow \text{dst} + \text{src}$

`subq src,dst`

$\text{dst} \leftarrow \text{dst} - \text{src}$

`incq dst`

$\text{dst} \leftarrow \text{dst} + 1$

`decq dst`

$\text{dst} \leftarrow \text{dst} - 1$

`negq dst`

$\text{dst} \leftarrow -\text{dst}$

(2's complement arithmetic negation)

x86-64 Memory Stack

- ▶ Register `%rsp` points to the “top” of stack
 - ▶ Dedicated for this use; don’t use otherwise
 - ▶ Points to the last 64-bit quadword pushed onto the stack (not next “free” quadword)
 - ▶ Should always be quadword (8-byte) aligned
 - ▶ It will start out this way, and will stay aligned unless your code does something bad
 - ▶ Should be 16-byte aligned on most function calls
 - ▶ Stack grows down

Stack Instructions

pushq src

$\text{\%rsp} \leftarrow \text{\%rsp} - 8;$
 $\text{memory}[\text{\%rsp}] \leftarrow \text{src}$
(e.g., push src onto the stack)

popq dst

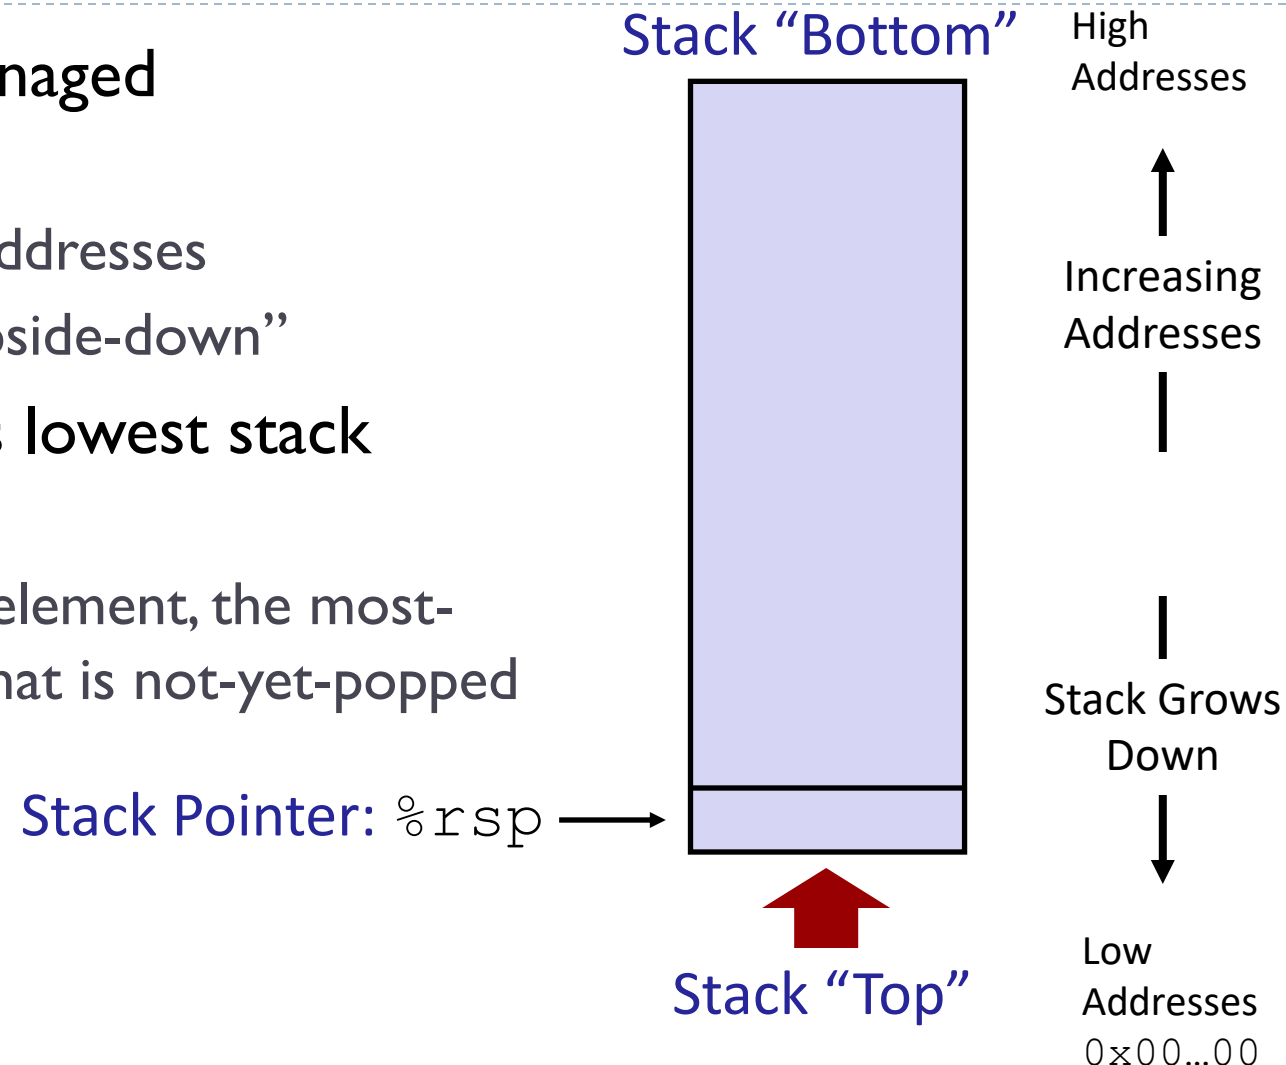
$\text{dst} \leftarrow \text{memory}[\text{\%rsp}];$
 $\text{\%rsp} \leftarrow \text{\%rsp} + 8$
(e.g., pop top of stack into dst and logically remove it from the stack)

Stack Frames

- ▶ When a method is called, a stack frame is traditionally allocated on the top of the stack to hold its local variables
- ▶ Frame is popped on method return
- ▶ By convention, `%rbp` (base pointer) points to a known offset into the stack frame
 - ▶ Local variables referenced relative to `%rbp`
 - ▶ Base pointer common in 32-bit x86 code; less so in x86-64 code where push/pop used less & stack frame normally has fixed size so locals can be referenced from `%rsp` easily
 - ▶ We will use `%rbp` in our project – simplifies addressing of local variables and compiler bookkeeping

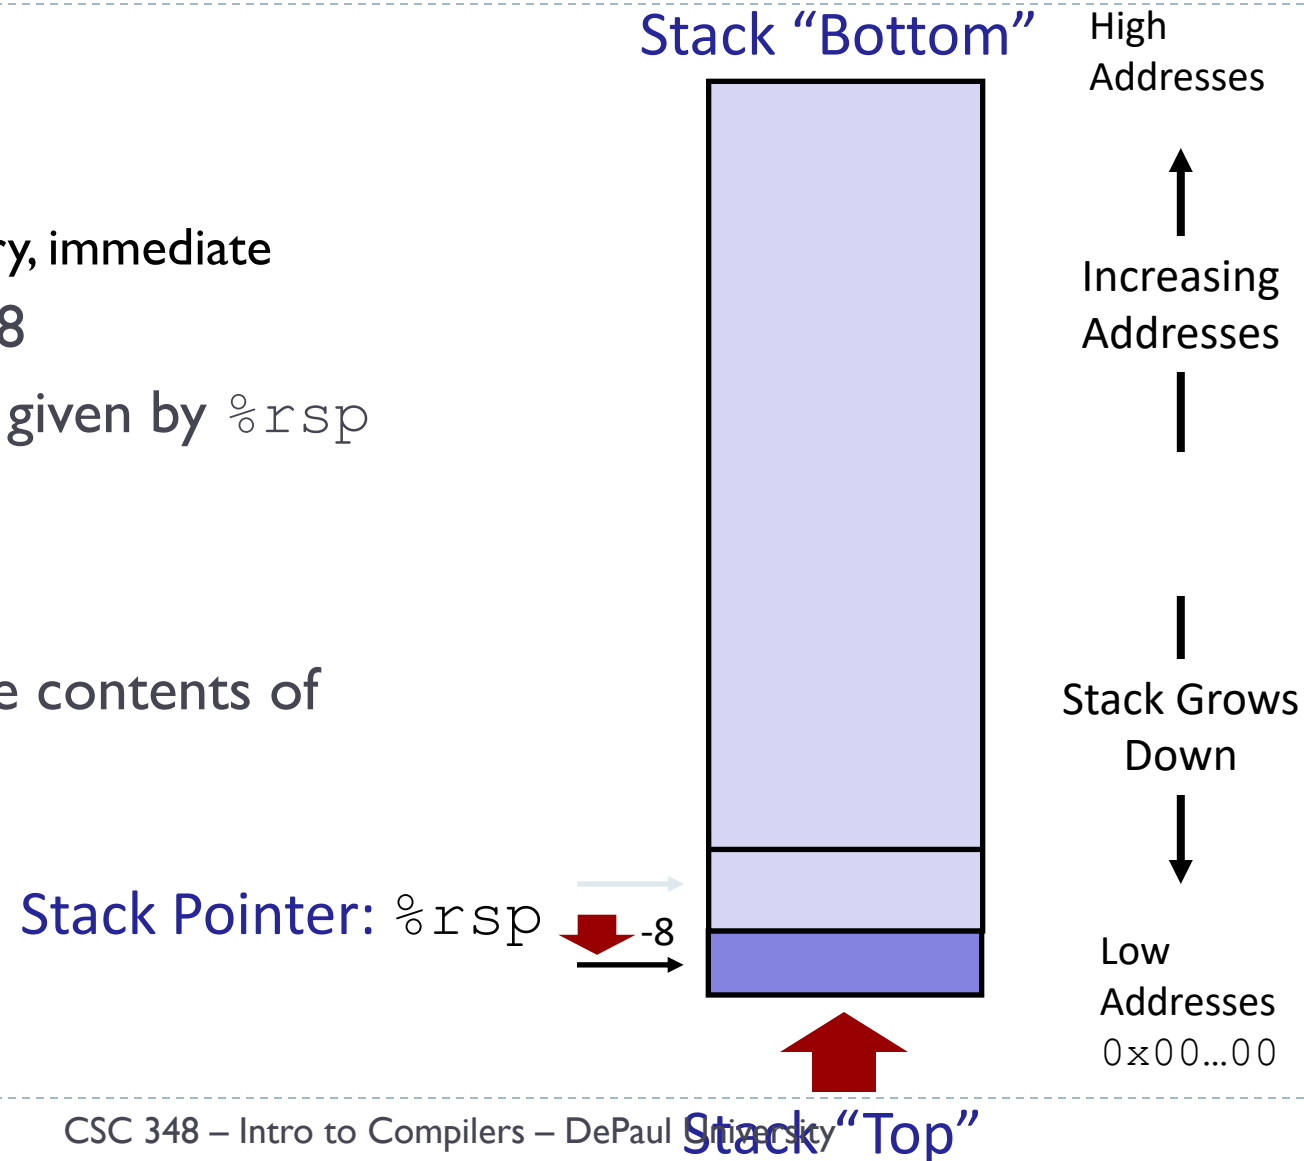
x86-64 Stack

- ▶ Region of memory managed with stack “discipline”
 - ▶ Grows toward lower addresses
 - ▶ Customarily shown “upside-down”
- ▶ Register `%rsp` contains lowest stack address
 - ▶ `%rsp` = address of top element, the most-recently-pushed item that is not-yet-popped



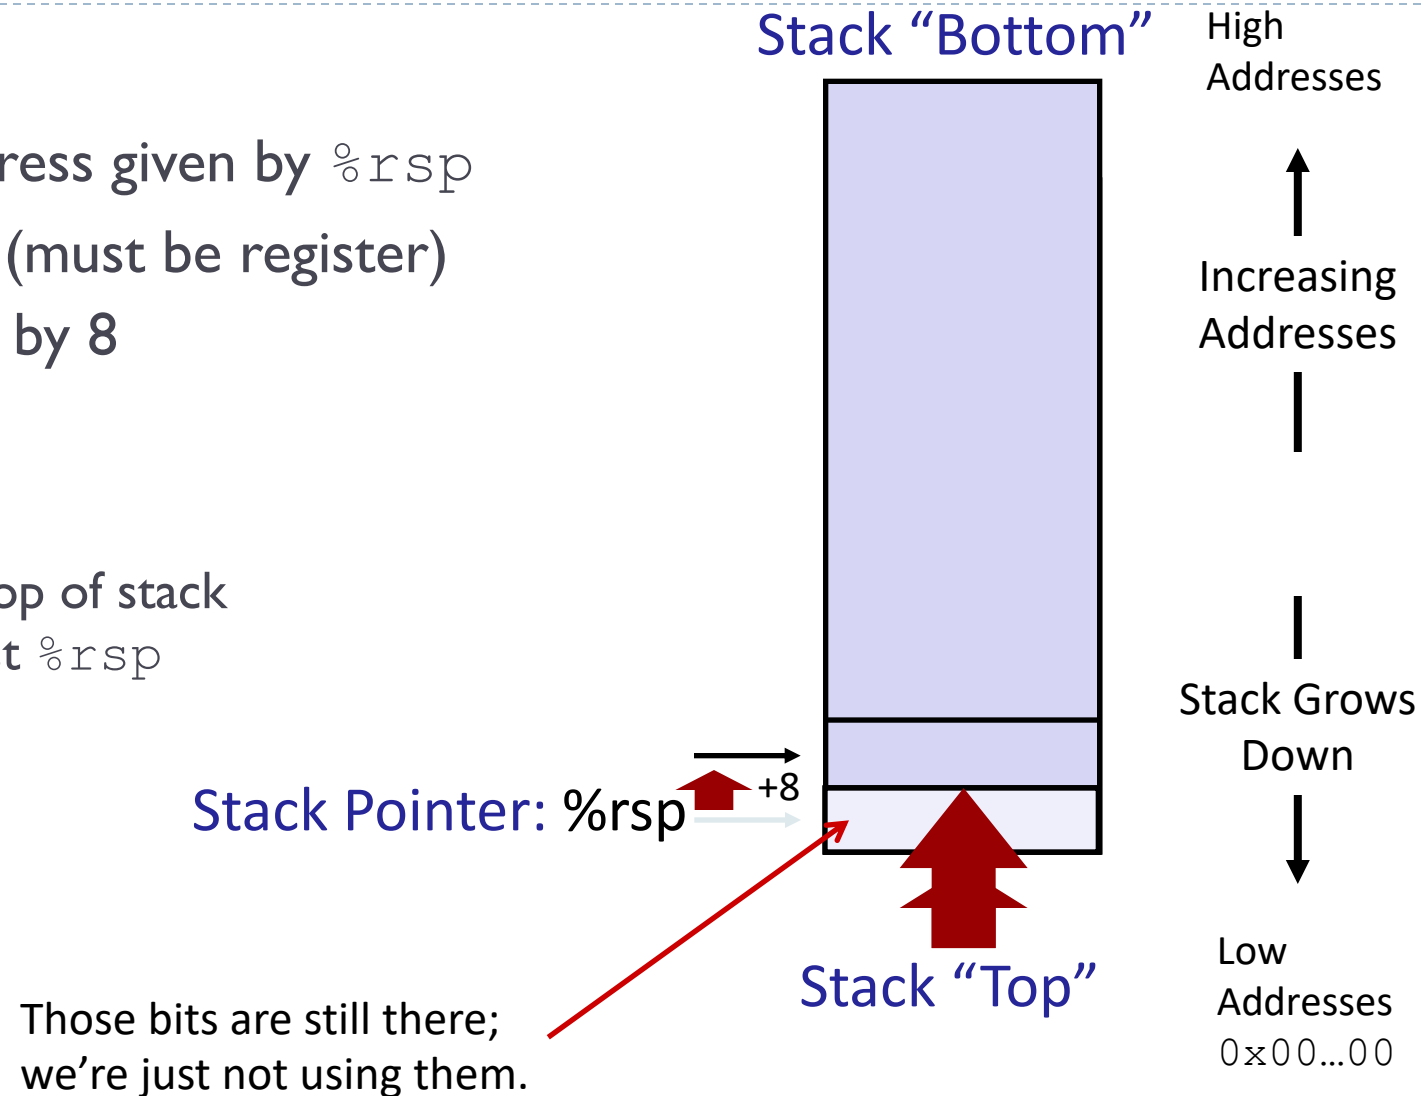
x86-64 Stack: Push

- ▶ `pushq src`
 - ▶ Fetch operand at `src`
 - ▶ `Src` can be reg, memory, immediate
 - ▶ **Decrement** `%rsp` by 8
 - ▶ Store value at address given by `%rsp`
- ▶ Example:
 - ▶ `pushq %rcx`
 - ▶ Adjust `%rsp` and store contents of `%rcx` on the stack



x86-64 Stack: Pop

- ▶ `popq dst`
 - ▶ Load value at address given by `%rsp`
 - ▶ Store value at *dst* (must be register)
 - ▶ **Increment** `%rsp` by 8
- ▶ Example:
 - ▶ `popq %rcx`
 - ▶ Stores contents of top of stack into `%rcx` and adjust `%rsp`



Operand Address Modes (1)

- ▶ These should cover most of what we'll need

<code>movq \$17,%rax</code>	<code># store 17 in %rax</code>
<code>movq %rcx,%rax</code>	<code># copy %rcx to %rax</code>
<code>movq 16(%rbp),%rax</code>	<code># copy memory to %rax</code>
<code>movq %rax,-24(%rbp)</code>	<code># copy %rax to memory</code>

- ▶ References to object fields work similarly – put the object's memory address in a register and use that address plus an offset
- ▶ Remember: can't have two memory addresses in a single instruction

Operand Address Modes (2)

- ▶ A memory address can combine the contents of two registers (with one optionally multiplied by 2, 4, or 8) plus a constant:

$\text{basereg} + \text{indexreg} * \text{scale} + \text{constant}$

- ▶ Main use of general form is for array subscripting or small computations - if the compiler is clever
- ▶ Example: suppose we have an array of 8-byte ints with address of the array A in %rcx and subscript i in %rax. Code to store %rbx in A[i]:

`movq %rbx, (%rcx, %rax, 8)`

Integer Multiply and Divide

`imulq src,dst`

$\text{dst} \leftarrow \text{dst} * \text{src}$

dst must be a register

`cqto`

$\%rdx:\%rax \leftarrow$ 128-bit sign extended copy of $\%rax$

(why??? To prep numerator for `idivq`!)

`idivq src`

Divide $\%rdx:\%rax$ by `src` ($\%rdx:\%rax$ holds sign-extended 128-bit value; cannot use other registers for division)

$\%rax \leftarrow$ quotient

$\%rdx \leftarrow$ remainder

(no division in MiniJava!)

Bitwise Operations

andq src,dst

$\text{dst} \leftarrow \text{dst} \& \text{src}$

orq src,dst

$\text{dst} \leftarrow \text{dst} | \text{src}$

xorq src,dst

$\text{dst} \leftarrow \text{dst} \wedge \text{src}$

notq dst

$\text{dst} \leftarrow \sim \text{dst}$

(logical or 1's complement)

Shifts and Rotates

shlq dst,count

dst \leftarrow dst shifted left count bits

shrq dst,count

dst \leftarrow dst shifted right count bits (0 fill)

sarq dst,count

dst \leftarrow dst shifted right count bits (sign bit fill)

rolq dst,count

dst \leftarrow dst rotated left count bits

rorq dst,count

dst \leftarrow dst rotated right count bits

Uses for Shifts and Rotates

- ▶ Can often be used to optimize multiplication and division by small constants
 - ▶ But be careful – be sure semantics are OK
 - ▶ Example: right shift is not the same as integer divide for negative numbers (why?)
- ▶ There are additional instructions that:
 - ▶ shift and rotate double words
 - ▶ use a calculated shift amount instead of a constant
 - ▶ etc.

Load Effective Address

- ▶ The unary & operator in C/C++

`leaq src,dst` $\# \text{ dst} \leftarrow \text{address of src}$

- ▶ `dst` must be a register
- ▶ Address of `src` includes any address arithmetic or indexing
- ▶ Useful to capture addresses for pointers, reference parameters, etc.
- ▶ Also useful for computing arithmetic expressions that match $r1 + \text{scale} * r2 + \text{const}$

Control Flow - GOTO

- ▶ At this level, all we have is Goto and Conditional Goto
- ▶ Loops and conditional statements are synthesized from these
- ▶ Note on Goto:
 - ▶ random jumps play havoc with pipeline efficiency
 - ▶ much work is done in modern compilers and processors to minimize this impact

Unconditional Jumps

`jmp dst`

`%rip` \leftarrow address of `dst`

- ▶ `dst` is usually a label in the code (which can be on a line by itself)
- ▶ `dst` address can also be indirect using the address in a register or memory location (`*reg` or `*(reg)`) – use for method calls, switch

Conditional Jumps

- ▶ Most arithmetic instructions set “condition code” bits to record information about the result (zero, non-zero, >0 , etc.)
 - ▶ Sets condition codes: `addq`, `subq`, `andq`, `orq`;
 - ▶ Does not set conditions codes: `imulq`, `idivq`, `leaq`
- ▶ Other instructions that set condition codes
 - `cmpq src,dst` # compare dst to src (e.g., `dst-src`)
 - `testq src,dst` # calculate `dst & src` (logical and)
 - ▶ These do not alter `src` or `dst`

Conditional Jumps Following Arithmetic Operations

- ▶ jz label # jump if result == 0
- ▶ jnz label # jump if result != 0
- ▶ jg label # jump if result > 0
- ▶ jng label # jump if result <= 0
- ▶ jge label # jump if result >= 0
- ▶ jnge label # jump if result < 0
- ▶ jl label # jump if result < 0
- ▶ jnl label # jump if result >= 0
- ▶ jle label # jump if result <= 0
- ▶ jnle label # jump if result > 0
- ▶ Obviously, the assembler is providing multiple opcode mnemonics for several actual instructions

Compare and Jump Conditionally

- ▶ We would like to be able to compare two operands and jump if a relationship holds between them
- ▶ Would like to do this
`jmpcond op1, op2, label`
- ▶ but can't, because 3-operand instructions can't be encoded in x86-64
- ▶ Many other architectures share this same limitation

cmp and jcc

- ▶ Instead, we use a 2-instruction sequence

```
    cmpq    op1,op2  
    jcc     label
```

- ▶ where j_{cc} is a conditional jump that is taken if the result of the comparison matches the condition cc

- ▶ `je label` # jump if $op1 == op2$
- ▶ `jne label` # jump if $op1 \neq op2$
- ▶ `jg label` # jump if $op1 > op2$
- ▶ `jng label` # jump if $op1 \leq op2$
- ▶ `jge label` # jump if $op1 \geq op2$
- ▶ `jl label` # jump if $op1 < op2$
- ▶ `jle label` # jump if $op1 \leq op2$

Function Call and Return

- ▶ The x86-64 instruction set itself only provides for transfer of control (jump) and return
- ▶ Stack is used to capture return address and recover it
- ▶ Everything else – parameter passing, stack frame organization, register usage – is a matter of convention and not defined by the hardware

call and ret Instructions

call label

- ▶ Push address of next instruction and jump
- ▶ $\text{\%rsp} \leftarrow \text{\%rsp} - 8$
- ▶ $\text{memory}[\text{\%rsp}] \leftarrow \text{\%rip}$
- ▶ $\text{\%rip} \leftarrow \text{address of label}$
- ▶ Address can also be in a register or memory as with jmp – we'll use these for dynamic dispatch of method calls (more later)

ret

- ▶ Pop address from top of stack and jump
- ▶ $\text{\%rip} \leftarrow \text{memory}[\text{\%rsp}]$;
- ▶ $\text{\%rsp} \leftarrow \text{\%rsp} + 8$
- ▶ **WARNING!** The word on the top of the stack had better be an address and not some leftover data

enter and leave

- ▶ Complex instructions for languages with nested procedures

- ▶ **enter** can be slow on current processors – best avoided – i.e., don't use it in your project

- ▶ **leave** is equivalent to

- `mov %rsp,%rbp`

- `pop %rbp`

and is generated by many compilers. Fits in 1 byte, saves space. Not clear if it's any faster.

X86-64-Register Usage for Functions

- ▶ **%rax** – function result
- ▶ Arguments 1-6 passed in these registers in order
 - ▶ **%rdi, %rsi, %rdx, %rcx, %r8, %r9**
 - ▶ For Java/C++ “this” pointer is first argument, in %rdi
- ▶ **%rsp** – stack pointer
 - ▶ value must be 8-byte aligned always and 16-byte aligned when calling a function
- ▶ **%rbp** – frame pointer (optional use)
 - ▶ We’ll use it

x86-64 Register Save Conventions

- ▶ A called function must preserve these registers (or save/restore them if it wants to use them)
 - ▶ `%rbx, %rbp, %r12 - %r15`
- ▶ `%rsp` isn't on the "callee save list", but needs to be properly restored for return
- ▶ All other registers can change across a function call
 - ▶ Debugging/correctness note: always assume every called function will change all registers it is allowed to

x86-64 Function Call

- ▶ Caller places up to 6 arguments in registers, rest on stack, then executes call instruction (which pushes 8-byte return address)
- ▶ On entry, called function prologue sets up the stack frame:

```
pushq    %rbp                # save old frame ptr
movq     %rsp,%rbp           # new frame ptr is top of stack after ret addr and
                              #   old rbp pushed
subq     $framesize,%rsp     # allocate stack frame
```

x86-64 Function Return

- ▶ Called function puts result (if any) in %rax and restores any callee-save registers if needed
- ▶ Called function returns with:

```
movq    %rbp,%rsp    # or use leave instead
popq    %rbp          #   of movq/popq
ret
```

- ▶ If caller allocated space for arguments it deallocates as needed

Caller Example

► `n = sumOf(17,42)`

<code>movq</code>	<code>\$17, %rdi</code>	<code># load arguments</code>
<code>movq</code>	<code>\$42, %rsi</code>	
<code>call</code>	<code>sumOf</code>	<code># jump & push ret addr</code>
<code>movq</code>	<code>%rax, offset_n(%rbp)</code>	<code># store result</code>

Example Function

► Source code

```
int sumOf(int x, int y) {  
    int a, int b;  
    a = x;  
    b = a + y;  
    return b;  
}
```


Assembly Language Version

```
# int sumOf(int x, int y) {  
# int a, int b;  
sumOf:  
    pushq    %rbp                # prologue  
    movq     %rsp, %rbp  
    subq     $16, %rsp  
  
# a = x;  
    movq     %rdi, -8(%rbp)      # %rdi = x  
  
# b = a + y;  
    movq     -8(%rbp), %rax  
    addq     %rsi, %rax          # %rsi = y  
    movq     %rax, -16(%rbp)  
  
# return b;  
    movq     -16(%rbp), %rax     # %rax = result  
    movq     %rbp, %rsp         # epilogue  
    popq     %rbp  
    ret  
# }
```

The Nice Thing About Standards...

- ▶ The above is the System V/AMD64 ABI convention (used by Linux, OS X)
- ▶ Microsoft's x64 calling conventions are slightly different (sigh...)
 - ▶ First four parameters in registers %rcx, %rdx, %r8, %r9; rest on the stack
 - ▶ Stack frame needs to include empty space for called function to save values passed in parameter registers if desired
- ▶ Not relevant for us, but worth being aware of it

Next...

- ▶ Now that we've got a basic idea of the x86-64 instruction set, we need to map language constructs to x86-64
- ▶ Next class we will talk about Code Shape
- ▶ Then need to figure out how to get compiler to generate this and how to bootstrap things to run the compiled programs
- ▶ Programming Assignment 4: Semantic Analysis