

CSC 373 Spring 2019 Prof. Lytinen

Introduction to Assembly Language

Readings:

Bryant and O'Hallaron, sections 3.1-3.7; 3.10.

Notes do not cover 3.6.8

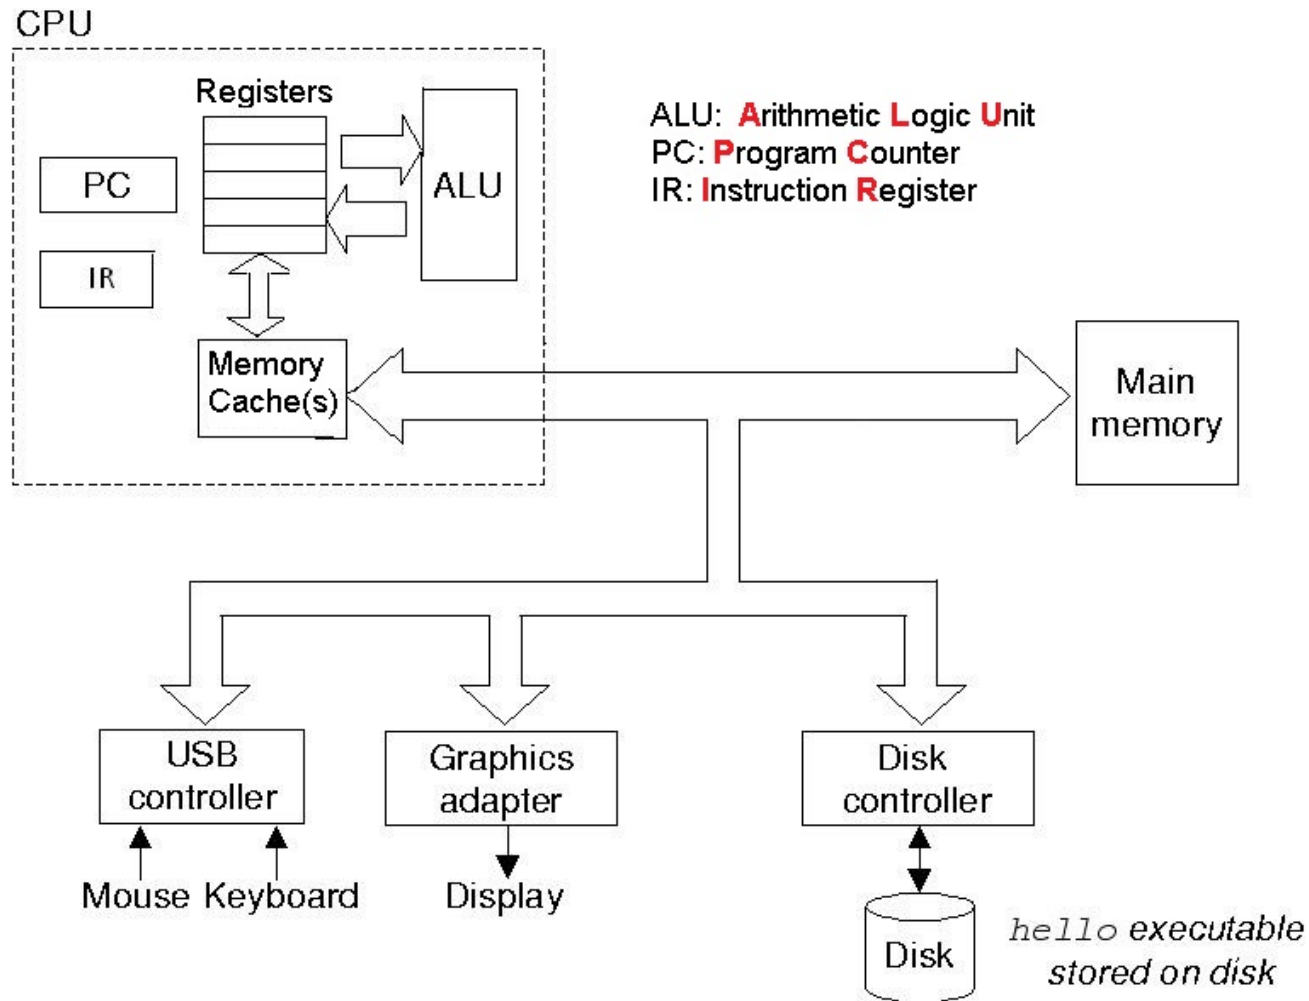
Machine Code

- When a C program is compiled, eventually it is translated into **machine code**
- Currently x86-64 machine code, named for historical reasons
- Original developed by **Advanced Micro Devices** (AMD) and was named AMD64
- x86-64 machine code has evolved from 16-bit processors to current 64-bit; backwards compatability has been maintained

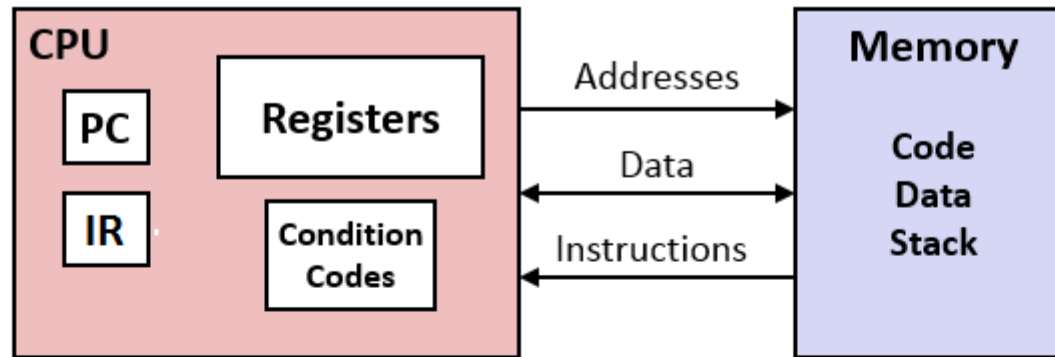
Assembly Language

- In essence, a (somewhat) more readable version of machine code
- Translation to machine code is almost one-to-one (i.e., each assembly language instruction often translates to one machine code instruction)
- Instructions' names are used; no need to memorize op code (see below)
- Likewise for some operands (although hexadecimal is used to indicate addresses)
- x64 Assembly

x86-64 Architecture



What your program sees



Programmer-Visible State

- **PC: Program counter**
 - Address of next instruction
 - Called "RIP" (x86-64)
- **Registers: 16 integer, 16 float point**
 - Heavily used program data
- **Condition codes**
 - Store status information about most recent arithmetic or logical operation

■ Memory

From your program's perspective, every 64-bit computer has:

2^{64} bytes ("virtual" memory) - no distinction between disk/RAM/caches (at least for the most part)

Memory from the point of view of a process

- Only 3 types of memory: **registers**, **main memory** and **files**
- Cache(s) are invisible to a process
- Processes also don't know what portions of their address space is currently in main memory and what is still on the disk.

Registers

- Program counter register (called %rip in assembly language) keeps track of execution location in main memory (a 64-bit address) . No direct access to this register.
- Instruction register contains the current instruction (1-15 bytes on x86-64)
- 16 regular 64-bit integer registers
- 16 64-bit floating point registers

Comparisons of speed

Memory Type	Latency (how many times slower)
Registers	1
L1 Cache	4
L2 Cache	10
L3 Cache	50
RAM	200
Disk	100,000

Chip history

- The first commercial 8-bit processor was the [Intel 8008](#) (1972) (Wikipedia)
- chars, ints, floats(?), addresses, etc. all took up 1 byte
- 2^8 bytes of addressable memory!! 0x00 – 0xff. Everything else had to be on a disc (or another computer)

Intel 8008: 8-bit machine

- Eight 8-bit registers: A, B, C, D, SI, DI, SP, BP
- First 4 were general purpose; the other 4 were for specific purposes
- Naming made some sense at the time (SI = Source Index, DI = Destination Index, SP = Stack pointer, BP = Base pointer)

Intel 8086: 16 bit machine

Intel 8086 registers

1 1 1 1 1 1 1 1 1 1 0 0 0 0 0 0 0 0 0 0 (bit position)
9 8 7 6 5 4 3 2 1 0 9 8 7 6 5 4 3 2 1 0

Main registers H = higher, L = lower

	AH	AL	AX (primary accumulator)
	BH	BL	BX (base, accumulator)
	CH	CL	CX (counter, accumulator)
	DH	DL	DX (accumulator, extended acc.)

Index registers

0 0 0 0	SI	Source Index	16 bits each
0 0 0 0	DI	Destination Index	
0 0 0 0	BP	Base Pointer	
0 0 0 0	SP	Stack Pointer	

Program counter

0 0 0 0	IP	Instruction Pointer
---------	----	---------------------

32-bit machines




Figure 1. x86 Registers

X86-64 Integer Registers

<code>%rax</code>	<code>%eax</code>
<code>%rbx</code>	<code>%ebx</code>
<code>%rcx</code>	<code>%ecx</code>
<code>%rdx</code>	<code>%edx</code>
<code>%rsi</code>	<code>%esi</code>
<code>%rdi</code>	<code>%edi</code>
<code>%rsp</code>	<code>%esp</code>
<code>%rbp</code>	<code>%ebp</code>

<code>%r8</code>	<code>%r8d</code>
<code>%r9</code>	<code>%r9d</code>
<code>%r10</code>	<code>%r10d</code>
<code>%r11</code>	<code>%r11d</code>
<code>%r12</code>	<code>%r12d</code>
<code>%r13</code>	<code>%r13d</code>
<code>%r14</code>	<code>%r14d</code>
<code>%r15</code>	<code>%r15d</code>

- Can reference low-order 4 bytes (also low-order 1 & 2 bytes)

Example: `%rdi`, `%edi`, `%di`, `%dil`

x86-64 instructions

- Machine code instructions are composed of an **operation code** (op code) and its **operands**. Of course, all are binary numbers.
- The first portion of an instruction (up to 3 bytes) encodes the op code
- Instructions are very low-level; e.g., 0x4889 is the op code which copies data from one 64-bit register to another

- Example:

48 89 d8 means `movq %rbx,%rax`

- Aside: why is the entire instruction 3 bytes long?
- Each opcode has a fixed number of operands (1 or 2), which takes up a fixed number of bytes (or possibly bits)
- Each operand specifies either a register or a (64-bit virtual) memory address
Registers can be encoded in 4 bits, since there are 16 of them.

Register names

- General-purpose registers are legacy names for backwards compatibility
- First microprocessors were 8-bit (meaning the address space of a process was 2^8 bytes)
- General purpose registers were %a, %b, %c, %d
- Other registers had special purposes, and their names reflect this for historical reasons. For example, only certain registers could have pointer values (%di, %si, %bp, %sp) which we'll talk about later
- si: source index; di = destination index; bp = base pointer; sp = stack pointer

Register names

- When 16-bit architecture was introduced, integer registers were increased to 2 bytes (as well as address spaces). To distinguish, full 16-byte registers were named %ax, %bx, %cx, and %dx. However, for backwards compatibility, the lower-order byte of the 4 general-purpose registers could still be used with the name %a, %b, %c, and %d.
- 32-bit architecture: %eax, %ebx, %ecx, %edx, %edi, %esi, %ebp, %esp
- 6 of the 8 registers became general-purpose, but retained their special-purpose names again for backwards compatibility (%ebp and %esp were still special-purpose)
- To reference 8 bits, the register names changed to %al, %ah, etc.
- 64-bit architecture: there are now 16 integer registers, called %rax, %rbx, %rcx, %rdx, %rdi, %rsi, %rbp, %rsp, %r8...%r15

X86-64 Integer Registers: Usage Conventions

While most registers can be used for general purposes, some are still also used for special purposes

Carne

%rax	Return value	%r8	Argument #5
%rbx	Callee saved	%r9	Argument #6
%rcx	Argument #4	%r10	Caller saved
%rdx	Argument #3	%r11	Caller Saved
%rsi	Argument #2	%r12	Callee saved
%rdi	Argument #1	%r13	Callee saved
%rsp	Stack pointer	%r14	Callee saved
%rbp	Callee saved	%r15	Callee saved

Example of assembly language and machine code

```
int main() {  
    printf("Hi\n");  
    printf("Bye\n");  
}
```

Memory address	Machine code of instruction in that address	Corresponding assembly language
400530	55	push %rbp
400531	48 89 e5	mov %rsp,%rbp
400534	bf e0 0f 40 00	mov \$0x4005e0, %edi
400539	e9 d2 fe ff ff	callq 400410 <puts@plt>
40053e	bf 3e 0f 40 00	mov \$0x4005e3, %edi
400543	e8 c8 fe ff ff	callq 400410 <puts@plt>
400548	5d	pop %rbp
400549	c3	retq

We can inspect further to see what the function does

- `%edi` is used to pass one of the parameters to a function (in this case `puts`)
- At 400534, the address 0x4005e0 is placed in `%edi`
- At 40053e, the address 0x4005e3 is placed in `%edi`

- In the debugger, we can type

(gdb) x/s 0x4005e0

0x4005e0: "Hi"

(gdb) x/s 0x4005e3

0x4005e3: "Bye"

- The op code `bf` is
“move immediate” (also
indicated by `$`

Memory address	Machine code	Corresponding assembly language
400530	55	push %rbp
400531	48 89 e5	mov %rsp,%rbp
400534	bf e0 05 40 00	mov \$0x4005e0 , %edi
400539	e9 d2 fe ff ff	callq 400410 <puts@plt>
40053e	bf e3 05 40 00	mov \$0x4005e3 , %edi
400543	e8 c8 fe ff ff	callq 400410 <puts@plt>
400548	5d	pop %rbp
400549	c3	retq

How a program runs

./hello

1. A **process** is created
2. An **address space** is created; always 2^{64} bytes on a 64-bit machine
The program now thinks it has 2^{64} bytes of memory
 - In reality, part of this address space may be on disk, in main memory, and/or in cache memory. Much of it will not be anywhere.
3. The starting address of the program in the address space is loaded into the CPU's **program counter** (%rip)
 - In the example above, 0x0000000000400530

How a program runs

4. Based on the contents of the PC, The first (next) instruction is copied from memory into the **instruction register**
 - First instruction: 0x55
5. The program counter is incremented by the appropriate number (1-15, depends on the length of the instruction)
 - After 1st instruction, increment %rip by 1
 - After 2nd instruction, increment %rip by 3
6. The instruction in the instruction register is executed. Some instructions modify the PC (to make loops, call other functions, etc.)

Repeat steps 4-6

PC

400530

IR

55

Memory address	Machine code	Corresponding assembly language
400530	55	push %rbp
400531	48 89 e5	mov %rsp,%rbp
400534	bf e0 05 40 00	mov \$0x4005e0 , %edi
400539	e9 d2 fe ff ff	callq 400410 <puts@plt>
40053e	bf e3 05 40 00	mov \$0x4005e3 , %edi
400543	e8 c8 fe ff ff	callq 400410 <puts@plt>
400548	5d	pop %rbp
400549	c3	retq

PC

400531

IR

48 89 e5

Memory address	Machine code	Corresponding assembly language
400530	55	push %rbp
400531	48 89 e5	mov %rsp,%rbp
400534	bf e0 05 40 00	mov \$0x4005e0 , %edi
400539	e9 d2 fe ff ff	callq 400410 <puts@plt>
40053e	bf e3 05 40 00	mov \$0x4005e3 , %edi
400543	e8 c8 fe ff ff	callq 400410 <puts@plt>
400548	5d	pop %rbp
400549	c3	retq

PC

400534

IR

bf e0 05 40 00

"little-endian"

Memory address	Machine code	Corresponding assembly language
400530	55	push %rbp
400531	48 89 e5	mov %rsp,%rbp
400534	bf e0 05 40 00	mov \$0x4005e0, %edi
400539	e9 d2 fe ff ff	callq 400410 <puts@plt>
40053e	bf e3 05 40 00	mov \$0x4005e3, %edi
400543	e8 c8 fe ff ff	callq 400410 <puts@plt>
400548	5d	pop %rbp
400549	c3	retq

%edi

e0 0e 5 40 00

PC

400539

IR

e9 d2 fe ff ff

Memory address	Machine code	Corresponding assembly language
400530	55	push %rbp
400531	48 89 e5	mov %rsp,%rbp
400534	bf e0 05 40 00	mov \$0x4005e0, %edi
400539	e9 d2 fe ff ff	callq 400410 <puts@plt>
40053e	bf e3 05 40 00	mov \$0x4005e3, %edi
400543	e8 c8 fe ff ff	callq 400410 <puts@plt>
400548	5d	pop %rbp
400549	c3	retq

%edi

e0 05 40 00

PC

400410 (start of
puts)

IR

???

Memory address	Machine code	Corresponding assembly language
400530	55	push %rbp
400531	48 89 e5	mov %rsp,%rbp
400534	bf e0 05 40 00	mov \$0x4005e0, %edi
400539	e9 d2 fe ff ff	callq 400410 <puts@plt>
40053e	bf e3 05 40 00	mov \$0x4005e3 , %edi
400543	e8 c8 fe ff ff	callq 400410 <puts@plt>
400548	5d	pop %rbp
400549	c3	retq

%edi

e0 05 40 00

PC

40053e

IR

bf e3 05 40 00

Memory address	Machine code	Corresponding assembly language
400530	55	push %rbp
400531	48 89 e5	mov %rsp,%rbp
400534	bf e0 05 40 00	mov \$0x4005e0, %edi
400539	e9 d2 fe ff ff	callq 400410 <puts@plt>
40053e	bf e3 05 40 00	mov \$0x4005e3, %edi
400543	e8 c8 fe ff ff	callq 400410 <puts@plt>
400548	5d	pop %rbp
400549	c3	retq

%edi

e3 05 40 00

PC

400549

IR

c3

Back to OS

Memory address	Machine code	Corresponding assembly language
400530	55	push %rbp
400531	48 89 e5	mov %rsp,%rbp
400534	bf e0 05 40 00	mov \$0x4005e0, %edi
400539	e9 d2 fe ff ff	callq 400410 <puts@plt>
40053e	bf e3 05 40 00	mov \$0x4005e3, %edi
400543	e8 c8 fe ff ff	callq 400410 <puts@plt>
400548	5d	pop %rbp
400549	c3	retq

%edi

e3 05 40 00

Data formats

In Intel/AMD-speak, integer types have different names than in C. Again, this is for historical reasons

C declaration	Data type	AMD64 suffix	Size (bytes)
char	Byte	b	1
short	Word	w	2
int	double word	l	4
long	quad word	q	8
All pointers	quad word	q	8

Examples:

movq %rax,%rbx means move (copy) 64 bits from %rax to %rbx

movl %eax,%ebx means move (copy) the least significant 32 bits from %rax to %rbx. The most significant 32 bits are not affected.

Assembly Language Example

C Code: Add two signed integers

```
long add(long x, long y) { return x + y; }
```

Assembly

```
movq %rsi, %rax # copy y into %rax  
addq %rdi, %rax # add x to %rax  
ret
```

By convention, the first 2 parameters are passed in %rdi and %rsi. The return value is passed in %rax.

Sort of like

```
%rax = %rsi  
%rax += %rdi  
return %rax
```

Common x86-64 Instructions

In 2-operand instructions, 1st operand is *source*, 2nd operand is *destination*. In 1-operand instructions, operand is *destination*. Many instructions have variants, such as `mov`, `movq`, `movl`, `movw`, `movb`. For most of the operations, both operands may be registers, or one may be a memory address and the other a register. Almost no instructions allow 2 memory addresses as their operands.

Instruction	# of operands	Meaning
<code>mov</code>	2	copy from src to dest
<code>add</code>	2	add src to dest; store result in dest
<code>sub</code>	2	subtract src from dest; store result in dest
<code>imul</code>	2	integer multiplication; store result in dest
<code>cmp</code>	2	Compare dest with src; flags are set accordingly
<code>inc</code>	1	increment dest
<code>dec</code>	1	decrement dest
<code>neg</code>	1	negate dest (e.g., 1 becomes -1)

Common x86-64 Instructions

Instruction	# of operands	Meaning
not	1	bitwise not dest
and	2	bitwise and dest and src; store result in dest
or	2	bitwise or dest and src; store result in dest
xor	2	bitwise xor dest and src; store result in dest
sal	2	shift (arithmetic) left. dest is shifted by src bits.
sar	2	shift (arithmetic) right. dest is shifted by src bits.
shr	2	shift (logical) right. dest is shifted by src bits.
lea	2	Load effective address. To be discussed later.
push	1	Push a value onto the program stack. To be discussed later
pop	1	Pop a value from the program stack. To be discussed later.

Special registers

- Some registers are still **only** used for special purposes. These are:

%rip: instruction register. Contains the instruction currently being executed.
Almost never seen in assembly language code.

%rsp: stack pointer. Almost always used in one particular way (more later)

%rbp: base pointer. Almost always used in one particular way (more later)

- Others may be used for general purposes, but by convention are used as follows:

%rdi, %rsi, %rdx, %rcx, %r8, %r9: used for parameter passing (although they may also be used for other purposes)

%rax: used to return a value from a function (although it might be used for other purposes)

Examples

```
int same(int x) {  
    return x; }
```

```
same:  movl %edi, %eax  
       ret
```

```
int add(int x, int y) {  
    int s = x + y;  
    return s; }
```

```
add:  movl %rsi, %eax  
      addl %rdi, %eax  
      ret
```

```
unsigned long times8(unsigned long x) {  
    return x * 8; }
```

```
times8:  
    movq %rdi, %rax  
    shlq $3, %rax  
    ret
```

Examples

```
int negative(int x) {  
    return -x; }
```

```
negative: movl %edi,%eax  
         negl %eax  
         ret
```

```
char abs (char x) {  
    if (x >= 0) return x;  
    else return -x; }
```

```
abs:  
    movl    %edi, %edx  
    movl    %edi, %eax  
    sarb    $7, %dl  
    xorl    %edx, %eax  
    subl    %edx, %eax  
    ret
```

abs(1)

abs:

```
movl    %edi, %edx
movl    %edi, %eax
sarb    $7, %dl      # why 7?
xorl    %edx, %eax
subl    %edx, %eax
```

ret

%rdi/%edi	%rdx/%edx	%rax%eax
0x01	0x01	
0x01	0x01	0x01
0x01	0x00	0x01
0x01	0x00	0x01
0x01	0x00	0x01

abs(1) is 1

abs(-3)

abs:

```
movl    %edi, %edx
movl    %edi, %eax
sarb    $7, %dl
xorl    %edx, %eax
subl    %edx, %eax
ret
```

#rsi / #esi	#rdx / #edx	#rax / %eax
0xfd	0xfd	
0xfd	0xfd	0xfd
0xfd	0xff	0xfd
0xfd	0xff	0x02
0xfd	0xff	0x03

Notes:

- -3_{10} is 0xfd (in 8-bit 2s complement)
- $0xff \wedge 0xfd$ is 0x02 (see next slide)
- $0x02 - 0xff$ is 0x03 (using 8-bit 2s complement)

$$\text{abs}(-3) = 3$$

abs:

```
movl    %edi, %edx
movl    %edi, %eax
sarb    $7, %dl
xorl    %edx, %eax
subl    %edx, %eax
ret
```

abs(-3)

%rdi

0xfd

%rdx

0xff

%rax

0xfd -> 0x02

11111111

^

11111101

=

00000010

abs:

```
movq %rdi,%rdx # %rdx = x
sarq $63,%rdx  # shift right - why 63 bits?
movq %rdx, %rax
xorq %rdi, %rax
subq %rdx, %rax
ret
```

abs(-3)

%rdi

0xfd

%rdx

0xff

%rax

0x03

$$2 - (-1) = 3$$

Condition codes

- The CPU maintains 1-bit registers, which store information about the results of the previous instruction. Usually the instructions are a variant of `cmp` but other instructions also set the condition codes, such as `add`, `subtract`, `shift`, etc. However, there are very few instructions which allow us to directly examine these flags

Code	Meaning
CF (carry flag)	The most recent operation generated a carry out of the most significant bit. Overflow for unsigned operations, including left logical shifts.
ZF (zero flag)	The most recent operation yielded 0. Note that <code>cmp</code> performs a subtraction, so ZF indicates the the numbers are equal
SF (sign flag)	The most recent operation yielded a negative number.
OF (overflow flag)	The most recent operation yielded a 2s complement overflow (from positive to negative or vice versa)

The `cmp` and `test` instructions

- Both compare numbers
- `cmp` subtracts one from the other
e.g., `cmp %edx, %eax` computes `%eax - %edx`
- `test` performs a bitwise-and
- yields different condition codes in certain cases
- Example: `%eax` is 8, `%edx` is 6

`cmpl %edx, %eax` sets CF and SF

`testl %edx, %eax` sets ZF

`test` is often used to see if a register contains 0 (ZF is set)

set instructions

- Usually follow a `cmp` or `test` instruction
- Sets *dest* based on condition code registers

Instruction	# of operands	Meaning
<code>sete</code>	1	<i>Dest</i> is set to ZF
<code>setne</code>	1	<i>Dest</i> is set to \sim ZF
<code>setge</code>	1	<i>Dest</i> is set to ZF \sim SF
<code>setg</code>	1	<i>Dest</i> is set to \sim ZF & \sim SF
<code>setle</code>	1	<i>Dest</i> is set to ZF SF
<code>setl</code>	1	<i>Dest</i> is set to \sim ZF & SF

Example

```
int less(x, y) {  
    return x < y;  
}
```

less:

xorl %eax, %eax	# %eax = 0
cmpl %esi, %edi	# compare %edi with %esi (computes %edi - %esi)
setl %al	# if negative (%edi < %esi) , set %al to 1
ret	

Some (hopefully) easy ones

f:

```
movl    %edi, %eax
shrl    $31, %eax
ret
```

g:

```
movl    %edi, %eax
andl    $1, %eax
xorl    $1, %eax
ret
```

h:

```
cmpl    %esi, %edi
setl    %cl
xorl    %eax, %eax
cmpl    %edx, %esi
setl    %al
andl    %ecx, %eax
ret
```

64-bit register	32-bit register	Arg #
%rdi	%edi	1
%rsi	%esi	2
%rdx	%edx	3

Some (hopefully) easy ones

We didn't get to h on 1/30

f:

```
movl    %edi, %eax
shrl    $31, %eax
ret
```

g:

```
movl    %edi, %eax
andl    $1, %eax
xorl    $1, %eax
ret
```

h:

```
cmpl    %esi, %edi
setl    %cl
xorl    %eax, %eax
cmpl    %edx, %esi
setl    %al
andl    %ecx, %eax
ret
```

```
Int f(int x) {
    return x < 0;
}
```

```
int g(int x)
    return x % 2 == 0;
}
```

64-bit register	32-bit register	Arg #
%rdi	%edi	1
%rsi	%esi	2
%rdx	%edx	3

The `jmp` instructions

- Behavior depends on the condition flags
- Every `jmp` you could possibly imagine

OP	Jump if ...	Flags
<code>jmp</code>	always	-
<code>je</code>	equal	ZF
<code>js</code>	Sign	SF
<code>jns</code>	Not sign	\sim SF
<code>jg</code>	greater	\sim (SF^OF)& \sim ZF
<code>jge</code>	greater or equal	\sim (SF^OF)

OP	Jump if ...	Flags
<code>jl</code>	Less	SF ^ OF
<code>jle</code>	Less or equal	(SF ^ OF) ZF
<code>ja</code>	Above	\sim CF & \sim ZF
<code>jae</code>	Above or equal	\sim CF
<code>jb</code>	Below	CF
<code>jbe</code>	Below or equal	CF ZF

The C goto statement

C

```
int greater(int x, int y) {  
    int diff = x - y;  
    if (diff > 0)  
        goto retx;  
    return y;  
retx: return x;  
}
```

64-bit register	32-bit register	Arg #
%rdi	%edi	1
%rsi	%esi	2
%rdx	%edx	3

x86-64

```
greater:  
    cmpl    %edi, %esi  
    jl      .L2  
    movl    %edi, %eax  
    jmp     .L4  
.L2:  
    movl    %esi, %eax  
.L4:  
    ret
```

The C goto statement

```
int power(int x, int y) {  
    int a = 1;  
    int i=0;  
start: if (i == y)  
        goto end;  
    a *= x;  
    i++;  
    goto start;  
end:   return a;  
}
```

```

int power(int x, int y) {
    int a = 1;
    int i=0;
L6: if (i == y)
        return a;
    a *= x;
    i++;
    goto L6;
}

```

```

power:
    xorl    %edx, %edx
    testl   %esi, %esi
    movl    $1, %eax
    je      .L5
.L6:
    addl    $1, %edx
    imull   %edi, %eax
    cmpl    %esi, %edx
    jne     .L6
.L5:
    ret

```

64-bit register	32-bit register	Arg #
%rdi	%edi	1
%rsi	%esi	2
%rdx	%edx	3

The C goto statement: factorial function

factorial:

```
    cmpl    $1, %edi  
    jle     .L4  
    addl    $1, %edi  
    movl    $2, %edx  
    movl    $1, %eax
```

.L3:

```
    imull   %edx, %eax  
    addl    $1, %edx  
    cmpl    %edi, %edx  
    jne     .L3  
    ret
```

.L4:

```
    movl    $1, %eax  
    ret
```

Flow of control

- No if...else or looping constructs in assembly language
- Instead, `set`, `test`, `cmp`, `and` `jmp` are used to create the same effects
- “Spaghetti code”

Another goto example in C

```
void print_equal(int x, int y) {  
    if (x != y)  
        printf("Not equal\n");  
    else printf("Equal\n");  
  
void print_equal(int x, int y) {  
    if (x != y) {  
        printf("Not equal\n");  
        goto end;  
    }  
    printf("Equal\n");  
end: return  
}
```

64-bit register	32-bit register	Arg #
%rdi	%edi	1
%rsi	%esi	2
%rdx	%edx	3

```
print_equal:  
    cmpl    %esi, %edi  
    je      .L2  
    movl    $.LC0, %edi  
    jmp     puts  
  
.L2:  
    movl    $.LC1, %edi  
    jmp     puts  
  
print_equal_goto:  
    cmpl    %esi, %edi  
    je      .L5  
    movl    $.LC0, %edi  
    jmp     puts  
  
.L5:  
    movl    $.LC1, %edi  
    jmp     puts
```

Memory addressing

- Finite number of registers
- Eventually, "main memory" must be used to store working data
 - **Virtual** address space
- Difficult to demonstrate with simple C programs, unless compilation is not optimized

```
int mult(int x, int y) {  
    return x * y;  
}
```

```
> gcc -o mult.s mult.c -S -O2
```

```
mult:  
    movl    %edi, %eax  
    imull   %esi, %eax  
    ret
```

Memory addressing

```
int mult(int x, int y) {  
    return x * y;  
}
```

```
> gcc -o mult.s mult.c -S
```

```
mult:
```

pushq	%rbp	# ignore
movq	%rsp, %rbp	# ignore
movl	%edi, -4(%rbp)	# int a = x (a is in memory)
movl	%esi, -8(%rbp)	# int b = y (b is in memory)
movl	-4(%rbp), %eax	# ans = a
imull	-8(%rbp), %eax	# ans *= b
popq	%rbp	# ignore
ret		# return ans

Operand types

- **Immediate:** constant value in decimal or Hex; number preceded by \$
- **Register:** starts with %
- Memory reference

several different ways to specify an operand's memory address

Most generally D(B,I,s)

D = *displacement*

B = *base register*

I = *index register*

s = *scale*

Memory Operand types

D = displacement

B = base register

I = index register

s = scale

- **Absolute:** give memory location D (rarely used)
- **Indirect:** specify a register; it contains the memory location (B)
- **Base + displacement:** specify a register, add a value to its address (like pointer arithmetic) D(B)
- **Indexed:** specify 2 registers, or 2 registers + a constant (B,I) or D(B,I)
- **Scaled indexed:** multiply by the scale s (B,I,s) or D(B,I,s)

Addressing type	Syntax	Example	Result
Immediate	number preceded by \$	movq \$10, %rax	%rax set to 10
Register	% before name	movq %rax,%rdx	the contents of the register %rax are copied into %rdx
Memory	D(B,I,S)		
<p>Memory operands Can have 4 components:</p> <ol style="list-style-type: none"> 1. Displacement 2. Base 3. Index 4. Scale <p>Base and index are registers Displacement and scale are integers.</p>	<p>Any may be left out, although the syntax varies a bit</p>	<p>movq %rax, (%rbx) movl 4(%rbx),%rax movl (%rbx,%rcx,4), %rdx</p>	<p>Assume %rax is x, %rbx is the pointer p, %rdx is y, and %rcx contains 5</p> <p>*p = x x = p[1] y = p[5]</p>

C examples

data is in memory, not a register **lea: Load Effective Address**

Loads a memory address but does not retrieve/store data from that address
leaq -4(%rbp), %rdx Store the address computed by subtracting 4 from the contents of %rbp

Can also be used for arithmetic computations

leaq (,%rdi,4), %rax # %rax = x * 4

does **not** access memory

Scale is dependent on datatype size

<u>Data type</u>	<u>C example</u>	<u>Instruction example</u>
int	int x = 3;	movq \$3,-4(%rbp)
int []	y[i] = 0;	movq \$0,-32(%rbp,%rdx,4)
int *	p = &x;	leaq -4(%rbp),%rdx

Exercise

$$D(B,I,S) = B + (S*I) + D$$

Assume the following values are stores in the following values and registers. Fill in the following tables.

Memory Location	Value
0x100	0x000000ff
0x104	0x000000ab
0x108	0x00000013
0x10c	0x00000011

Operand	Value
%rax	
\$0x108	
(%rax)	
4(%rax)	

Register	Value
%rax	0x0000000000000100
%rcx	0x0000000000000001
%rdx	0x0000000000000003

Operand	Value
9(%rax, %rdx)	
256(%rcx,%edx)	
0xfc(,%ecx,4)	
(%rax,%rdx,4)	

Fill in the following table showing the effect of each of the instructions below. Assume the values in in memory and registers are as specified above. Assume the instructions are **not** sequential.

Instruction	Destination	New Value in Destination
movq %rax, (%rax)		
addl 4(%rax), %ecx		
subq %rdx, (%rax, %rcx, 4)		
movq \$-1, 4(%rax)		
movzbq \$0x61, 4(%rax,%rcx,4)		
movsbq \$-1,%rdx		

Note: the numbers in red should all be 64 bits, but for brevity I have written most of them as 16 bit numbers. This is also true of the previous slide.

Answers to Exercise $D(B,I,S) = B + (S*I) + D$

Assume the following values are stores in the following values and registers. Fill in the following tables.

Memory Location	Value
0x100	0x000000ff
0x104	0x000000ab
0x108	0x00000013
0x10c	0x00000011

Operand	Value
%rax	0x100
\$0x108	0x108
(%rax)	0xff
4(%rax)	0xab

Register	Value
%rax	0x0000000000000100
%rcx	0x0000000000000001
%rdx	0x0000000000000003

Operand	Value
9(%rax, %rdx)	0x11
256(%rcx,%edx)	0xab
0xfc(,%ecx,4)	0xff
(%rax,%rdx,4)	0x11

Answers to exercise

Fill in the following table showing the effect of each of the instructions below. Assume the values in in memory and registers are as specified above. Assume the instructions are **not** sequential.

Instruction	Destination	New Value in Destination
movq %rax, (%rax)	0x100	0x100
addl 4(%rax), %ecx	%ecx	0xac
subq %rdx, (%rax, %rcx, 4)	0x104	0xa8
movq \$-1, 4(%rax)	0x104	0xffffffffffffffff
movzbq \$0x61, 4(%rax,%rcx,4)	0x108	0x00000...00000061
movsbq \$-1,%rdx	%rdx	0xffffffffffffffff

Note: the numbers in red should all be 64 bits, but for brevity I have written most of them as 16 bit numbers. This is also true of the previous slide.

What does this function do?

f:

```
movl    (%rdi), %eax
movl    (%rsi), %edx
movl    %edx, (%rdi)
movl    %eax, (%rsi)
ret
```

64-bit register	32-bit register	Arg #
%rdi	%edi	1
%rsi	%esi	2
%rdx	%edx	3

lea – load effective address

Source: Memory operand $D(B, I, s)$

Value: $D(B, I, s)$ NOT THE CONTENTS OF
 $D(B, I, s)$

Like & operator in C

Sometimes used to do arithmetic

Fill in the following table showing the effect of each of the instructions below. Assume the values in in memory and registers are as specified above. Assume the instructions are **not** sequential.

Instruction	Destination	New Value in Destination
leal (%rax%, %rax), %rdx	%rdx	0x200
leal 4(%rax), %ecx	%eax	0x104
leal (%rcx, %rdx), %rax	%rax	0x4

Note: the numbers in red should all be 64 bits, but for brevity I have written most of them as 16 bit numbers. This is also true of the previous slide.

Strings as parameters

```
int mystrlen(char *s) {  
    int i;  
    for (i=0; s[i]!='\0'; i++);  
    return i;  
}
```

mystrlen:

```
    cmpb    $0, (%rdi)  
    je      .L4  
    addq    $1, %rdi  
    xorl    %eax, %eax
```

.L3:

```
    addq    $1, %rdi  
    addl    $1, %eax  
    cmpb    $0, -1(%rdi)  
    jne     .L3  
    ret
```

.L4:

```
    xorl    %eax, %eax  
    ret
```

Array vs. Pointer syntax example

```
int strlen373(char s[]) {  
    int i;  
    for (i=0; s[i] != '\0'; i++);  
    return i;  
}
```

```
int strlen373ptr(char *s) {  
    int i;  
    for (i=0; *s++ != '\0'; i++);  
    return i;  
}
```

```
strlen373:  
    cmpb    $0, (%rdi)  
    je      .L4  
    addq    $1, %rdi  
    xorl    %eax, %eax  
.L3:  
    addq    $1, %rdi  
    addl    $1, %eax  
    cmpb    $0, -1(%rdi)  
    jne     .L3  
    ret  
.L4:  
    xorl    %eax, %eax  
    ret
```

```
strlen373ptr:  
    xorl    %eax, %eax  
    cmpb    $0, (%rdi)  
    leaq    1(%rdi), %rdx  
    je      .L9  
.L8:  
    addq    $1, %rdx  
    addl    $1, %eax  
    cmpb    $0, -1(%rdx)  
    jne     .L8  
    ret  
.L9:  
    ret
```

64-bit register	32-bit register	Arg #
%rdi	%edi	1
%rsi	%esi	2
%rdx	%edx	3

```
int sum_of_squares (int n) {
    int ans = 0;
    int i;
    for (i=1; i<=n; i++)
        ans += i*i;
    return ans;
}
```

```
sum_of_squares:
    testl    %edi, %edi
    jle      .L4
    addl     $1, %edi
    movl     $1, %edx
    xorl     %eax, %eax
.L3:
    movl     %edx, %ecx
    imull    %edx, %ecx
    addl     $1, %edx
    addl     %ecx, %eax
    cmpl     %edi, %edx
    jne      .L3
    rep ret
.L4:
    xorl     %eax, %eax
    ret
```

```
int sum_of_squares(int n) {
    int ans = 0;
    int i=1;
    L1: if (i>n) goto L2;
        ans += i*i;
        i++;
        goto L1;
    L2: return ans;
}
```

64-bit register	32-bit register	Arg #
%rdi	%edi	1
%rsi	%esi	2
%rdx	%edx	3

The C code compiles the same in either case

```
// convert 'F' to 'C' and vice versa
void scale(char s1, char *s2) {
    if (s1 == 'F')
        *s2 = 'C';
    else if (s1 == 'C')
        *s2 = 'F';
}
```

scale:

```
    cmpb $70, %dil    # s1 == 'F'
    je .L6            # Yes? Goto L6
    cmpb $67, %dil    # s1 == 'C'
    je .L7            # Yes? Goto L7
    ret
.L7: movb $70, (%rsi)  # *s2 = 'F'; indirect addressing
    ret
.L6: movb $67, (%rsi)  # *s2 = 'C'; indirect addressing
    ret
```

64-bit register	32-bit register	8 bits	Arg #
%rdi	%edi	%dil	1
%rsi	%esi	%sil	2
%rdx	%edx	%dl	3

What does this function do?

f:

```
testl    %esi, %esi
jle      .L4
xorl     %edx, %edx
xorl     %eax, %eax
```

.L3:

```
addl     (%rdi,%rdx,4), %eax
addq     $1, %rdx
cmpl     %edx, %esi
jg       .L3
ret
```

.L4:

```
xorl     %eax, %eax
ret
```

64-bit register	32-bit register	Arg #
%rdi	%edi	1
%rsi	%esi	2
%rdx	%edx	3

- L3: memory operand base, index, scale
- How many parameters?
- What type of data?
- Is there a loop?

What does this function do?

```
f:
    testl    %esi, %esi    ← parameter
    jle      .L4
    xorl     %edx, %edx    ← Not parameter
    xorl     %eax, %eax

.L3:
    addl     (%rdi,%rdx,4), %eax
    addq     $1, %rdx
    cmpl     %edx, %esi    ← parameter
    jg       .L3
    ret

.L4:
    xorl     %eax, %eax
    ret
```

64-bit register	32-bit register	Arg #
%rdi	%edi	1
%rsi	%esi	2
%rdx	%edx	3

- L3: memory operand base, index, scale
- How many parameters? **2**
- Is there a loop?

What types of parameters?

Yes beginning with L3

%rdi: pointer

%rsi: integer


Returns the sum of an array of integers

Reverse Engineering problems

- Please see the accompanying .s file.

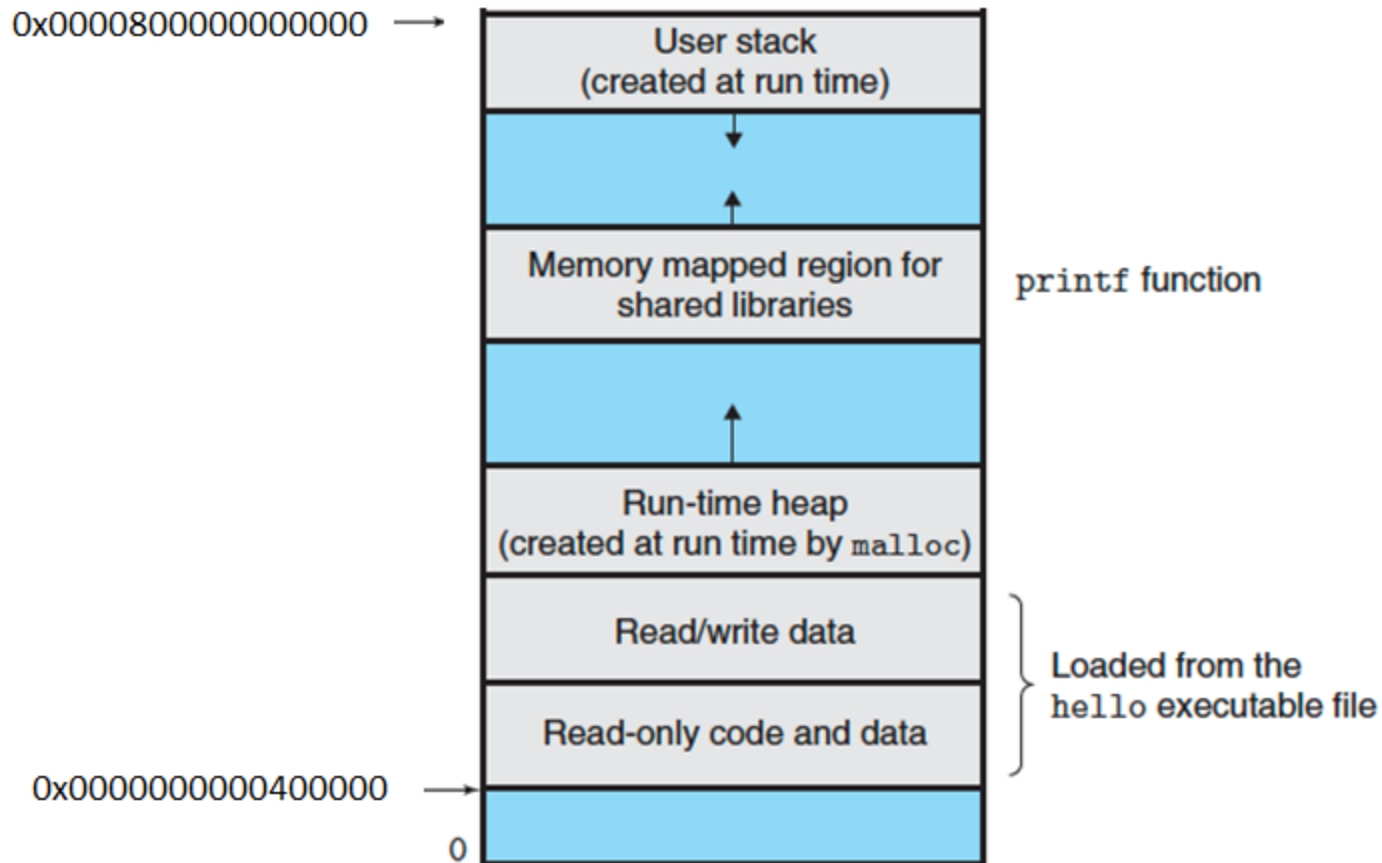
Division of address space

- Source: <http://thinkiii.blogspot.com/2014/02/arm64-linux-kernel-virtual-address-space.html>
- Not all 2^{64} bytes of the address space of a process are used;
- The “top” portion of the address space contains the OS “kernel” (its essentials); the bottom portion is your program and data. Note that the yellow is 2^{48} bytes, 2^{18} bytes are still unused.



User address space

<http://stackoverflow.com/questions/9511982/virtual-address-space-in-the-context-of-programming>



Program Stack

- Programs keep track of many things in a portion of memory called the **program stack**
- Instructions **push** and **pop** do what they sound like
- **%rsp** contains the address of the top of the stack
- When a function is called, it sets up a **stack frame** for itself
 - **%rsp** is a pointer to the top of the stack frame
 - **%rbp** is a pointer to the bottom of the stack frame

memory
addresses
increase



0x0000000000000000

Program
Stack

0xffffffffffffffff



Dynamically
grows (and
shrinks)



Address space
(2^{64} bytes)

Or if you prefer

memory
addresses
increase



0x0000000000000000

Program
Stack



↓
Dynamically
Grows (and
shrinks)

0xffffffffffffffff

Address space
(2^{64} bytes)

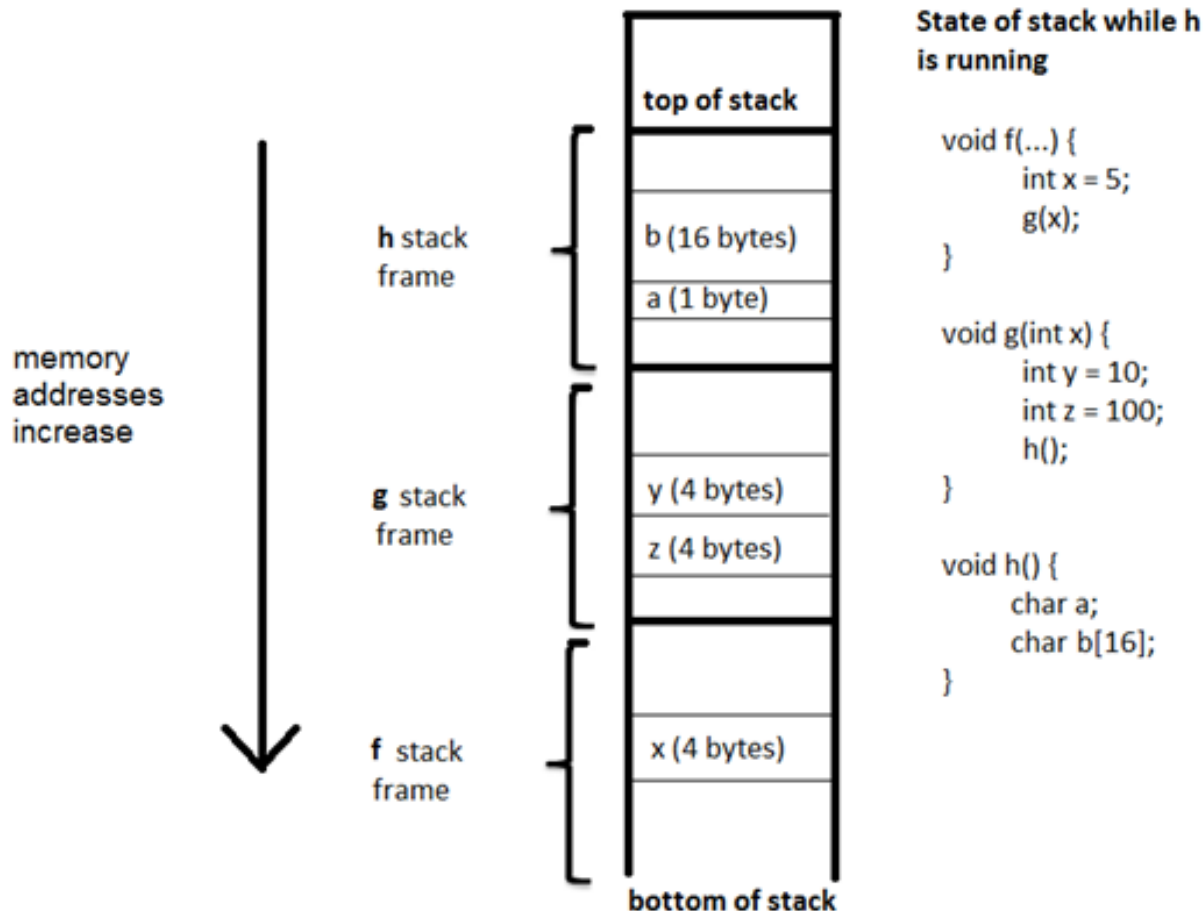


Orientation of the Program Stack

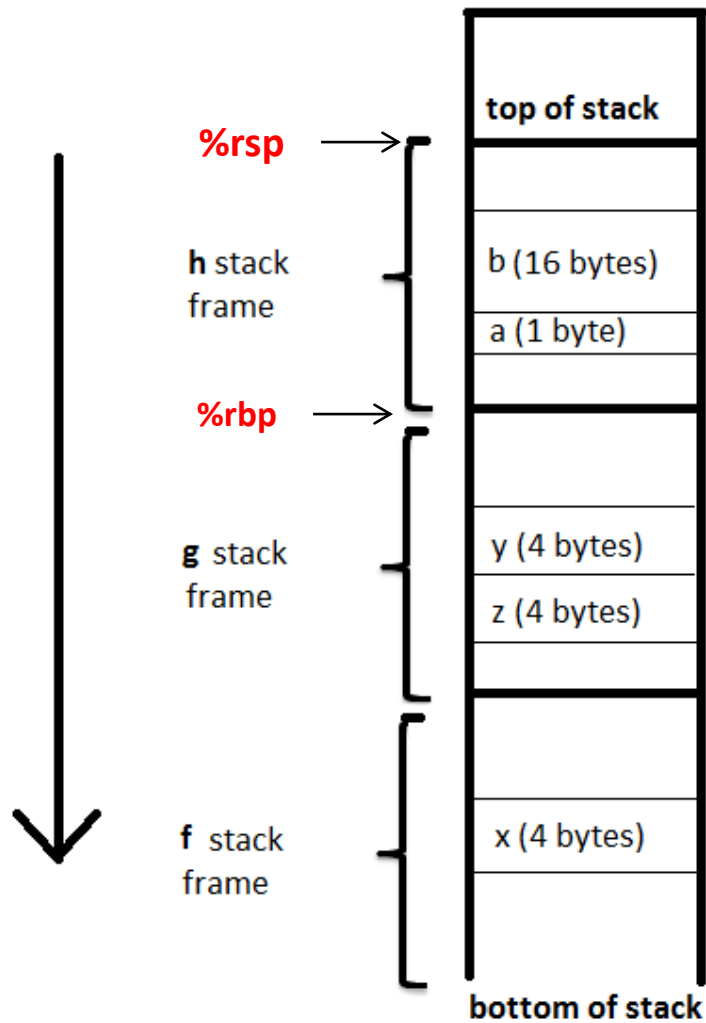
- Text. P. 90: “By convention, we draw stacks upside down, so that the “top” of the stack is shown at the bottom.”
- My diagrams **do not** follow this convention, because it doesn’t make sense.

Use of the program stack

- When a function is called, it allocates room on the stack for its “stack frame”
- The frame contains space for local variables, register values that need to be stored away, and information needed to return to the calling function




memory
addresses
increase



State of stack while **h**
is running

```
void f(...) {  
    int x = 5;  
    g(x);  
}  
  
void g(int x) {  
    int y = 10;  
    int z = 100;  
    h();  
}  
  
void h() {  
    char a;  
    char b[16];  
}
```

- Stack frame organization



```
void f() {  
    int x = 10;  
    g(x);  
}  
  
void g(int x) {  
    int y = 10;  
    int z = 100;  
    h();  
}
```

Register %rsp
contains the address
of the "top" of the
current stack frame

Register %rbp
contains the address
of the "bottom" of the
current stack frame

pushq

- When a function is called, it creates its own stack frame
- It also needs to store information that is required to return to the calling function
 - Return address
 - Boundaries of the calling function's stack frame
- Therefore, many functions start with

```
pushq %rbp  
movq %rsp, %rbp
```

- push is a “macro”, short for

```
subq $16, %rsp  
movq %rbp, (%rsp)
```

popq

- When a function returns, it restores the program stack to its previous state
- Therefore, many functions end with with

```
popq %rbp
```

- pop is a “macro”, short for

```
movq (%rsp), %rbp  
addq $16, %rsp
```


Example

```
int f(x) {  
    return g(x);  
}
```

```
int g(y) {  
    return y * 2;  
}
```

g:

```
pushq    %rbp  
movq      %rsp, %rbp  
subq      $-4, %rsp  
movl      %edi, -4(%rbp)  
movl      -4(%rbp), %eax  
addl      %eax, %eax  
addq      $4, %rsp  
popq      %rbp  
ret
```




Example

```
int f(x) {  
    return g(x);  
}
```

```
int g(y) {  
    return y * 2;  
}
```

%rbp → %rsp →



g:

```
pushq    %rbp  
movq    %rsp, %rbp  
subq     $-4, %rsp  
movl     %edi, -4(%rbp)  
movl     -4(%rbp), %eax  
addl     %eax, %eax  
addq     $4, %rsp  
popq     %rbp  
ret
```


Example

```
int f(x) {  
    return g(x);  
}
```

```
int g(y) {  
    return y * 2;  
}
```

g:

```
pushq    %rbp  
movq     %rsp, %rbp  
subq    $-4, %rsp  
movl     %edi, -4(%rbp)  
movl     -4(%rbp), %eax  
addl     %eax, %eax  
addq     $4, %rsp  
popq     %rbp  
ret
```




Example

```
int f(x) {  
    return g(x);  
}
```

```
int g(y) {  
    return y * 2;  
}
```

%rsp → %rbp →



g:

```
pushq    %rbp  
movq     %rsp, %rbp  
subq     $-4, %rsp  
movl     %edi, -4(%rbp)  
movl     -4(%rbp), %eax  
addl     %eax, %eax  
addq     $4, %rsp  
popq     %rbp  
ret
```


Example

```
int f(x) {  
    return g(x);  
}
```

```
int g(y) {  
    return y * 2;  
}
```

g:

```
pushq    %rbp  
movq     %rsp, %rbp  
subq     $-4, %rsp  
movl     %edi, -4(%rbp)  
movl     -4(%rbp), %eax  
addl     %eax, %eax  
addq     $4, %rsp  
popq     %rbp  
ret
```



pop example

less:

```
pushq  %rbp           # save %rbp from previous stack frame
movq   %rsp, %rbp      # %rbp now points to the bottom of "less" stack frame
subq   $16, %rsp       # adjust %rsp so that the stack frame is the right size
movl    %edi, -4(%rbp)
movl    %esi, -8(%rbp)
movl    -4(%rbp), %eax
cmpl    -8(%rbp), %eax
setl    %al
movzbl  %al, %eax
addq   $16, %rsp       # make "less" stack frame disappear
popq   %rbp           # restore calling function's base pointer
ret
```

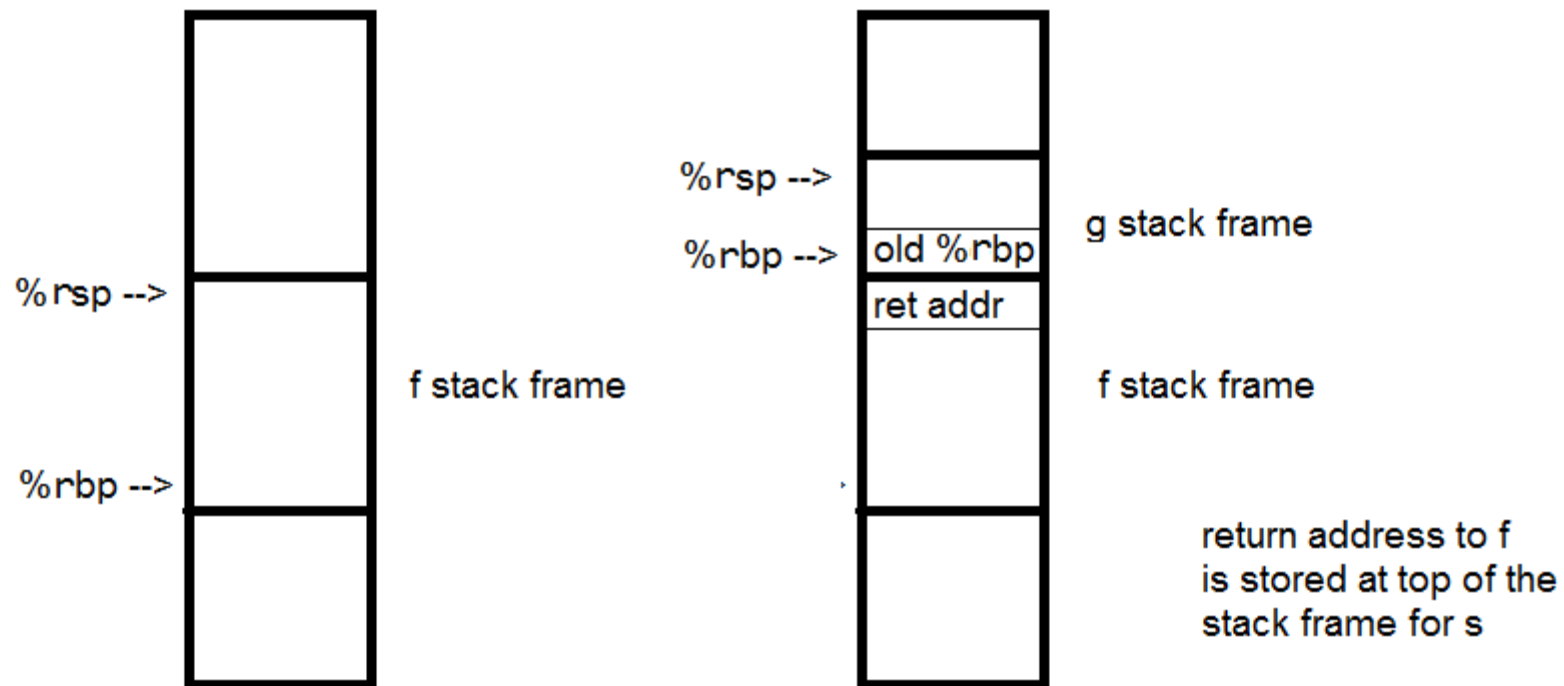

- **call and ret**

- Also macros; `callq 0x400f01` means

```
pushq %rip  
jmpq 0x400f01
```

and `ret` means

```
popq %rip  
jmpq *%rip # this is an indirect jump
```



Recursion

Recursive code is sometimes slower than iterative, because of the need to use the program stack.

Example: factorial

```
int fact_r(int x) {  
    if (x <= 1)  
        return 1;  
    else return x * fact_r(x-1);  
}
```

Where is x remembered?

First, using

`gcc -o fact.s -s fact.c`

fact_r:

```
    pushq %rbp
    movq  %rsp, %rbp
    subq  $16, %rsp
    movl  %edi, -4(%rbp)
    cmpl  $1, -4(%rbp)
    jg    .L2
    movl  $1, %eax
    jmp   .L3
```

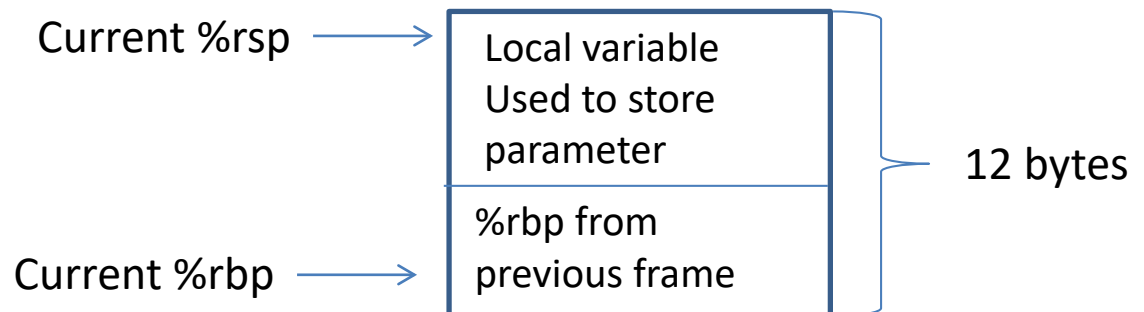
.L2:

```
    movl  -4(%rbp), %eax
    subl  $1, %eax
    movl  %eax, %edi
    call  fact_r
    imull -4(%rbp), %eax
```

.L3:

```
    leave
    ret
```

Each call to fact_r requires a stack frame



fact_r:

```
pushq %rbp
movq  %rsp, %rbp
subq  $16, %rsp
movl  %edi, -4(%rbp)
cmpl  $1, -4(%rbp)
jg    .L2
movl  $1, %eax
jmp   .L3
```

.L2:

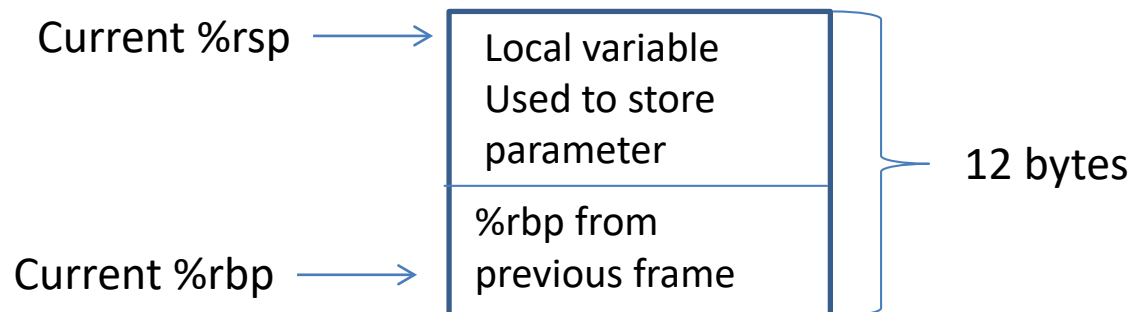
```
movl  -4(%rbp), %eax
subl  $1, %eax
movl  %eax, %edi
call  fact_r
imull -4(%rbp), %eax
```

.L3:

```
leave
ret
```

Each call to fact_r requires a stack frame

Establish stack frame



fact_r:

```
    pushq %rbp
    movq  %rsp, %rbp
    subq  $16, %rsp
    movl  %edi, -4(%rbp)
    cmpl  $1, -4(%rbp)
    jg    .L2
    movl  $1, %eax
    jmp   .L3
```

.L2:


```
    movl  -4(%rbp), %eax
    subl  $1, %eax
    movl  %eax, %edi
    call  fact_r
    imull -4(%rbp), %eax
```

.L3:

```
    leave
    ret
```

Each call to fact_r requires a stack frame

Store x



fact_r:

```
    pushq %rbp
    movq  %rsp, %rbp
    subq  $16, %rsp
    movl  %edi, -4(%rbp)
    cmpl  $1, -4(%rbp)
        jg  .L2
        movl $1, %eax
        jmp .L3
```

.L2:

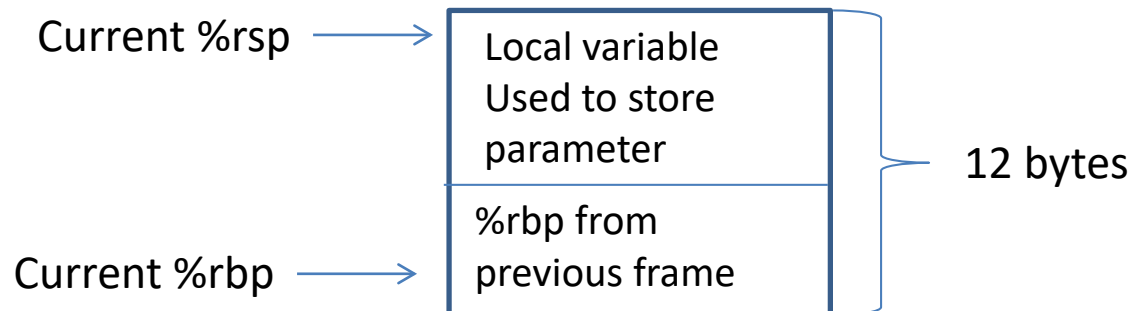
```
    movl  -4(%rbp), %eax
    subl  $1, %eax
    movl  %eax, %edi
    call  fact_r
    imull -4(%rbp), %eax
```

.L3:

```
    leave
    ret
```

Each call to fact_r requires a stack frame

(1)! Is 1



fact_r:

```
pushq %rbp
movq %rsp, %rbp
subq $16, %rsp
movl %edi, -4(%rbp)
cmpl $1, -4(%rbp)
jg .L2
movl $1, %eax
jmp .L3
```

.L2:


```
movl -4(%rbp), %eax
subl $1, %eax
movl %eax, %edi
call fact_r
imull -4(%rbp), %eax
```

.L3:

```
leave
ret
```

Each call to fact_r requires a stack frame

$(n)! \text{ is } (n)[(n-1)!]$



fact_r:

```
    pushq %rbp
    movq  %rsp, %rbp
    subq  $16, %rsp
    movl  %edi, -4(%rbp)
    cmpl  $1, -4(%rbp)
    jg    .L2
    movl  $1, %eax
    jmp   .L3
```


.L2:

```
    movl  -4(%rbp), %eax
    subl  $1, %eax
    movl  %eax, %edi
    call  fact_r
    imull -4(%rbp), %eax
```

.L3:

```
    leave
    ret
```

Consider fact_r(4):



fact_r:

```
pushq %rbp
movq  %rsp, %rbp
subq  $16, %rsp
movl  %edi, -4(%rbp)
cmpl  $1, -4(%rbp)
jg    .L2
movl  $1, %eax
jmp   .L3
```

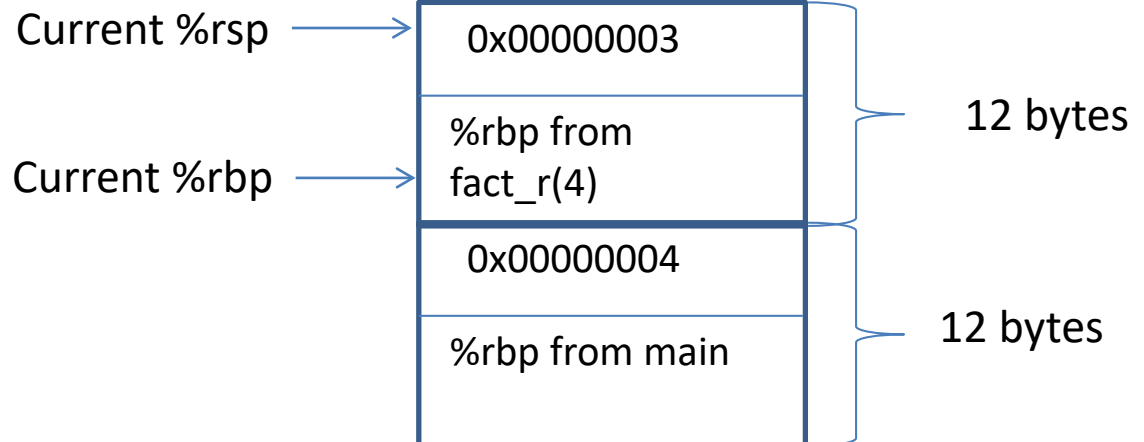
.L2:

```
movl  -4(%rbp), %eax
subl  $1, %eax
movl  %eax, %edi
call  fact_r
imull -4(%rbp), %eax
```

.L3:

```
leave
ret
```

fact_r(3):



fact_r:

```
pushq %rbp
movq  %rsp, %rbp
subq  $16, %rsp
movl  %edi, -4(%rbp)
cmpl  $1, -4(%rbp)
jg    .L2
movl  $1, %eax
jmp   .L3
```

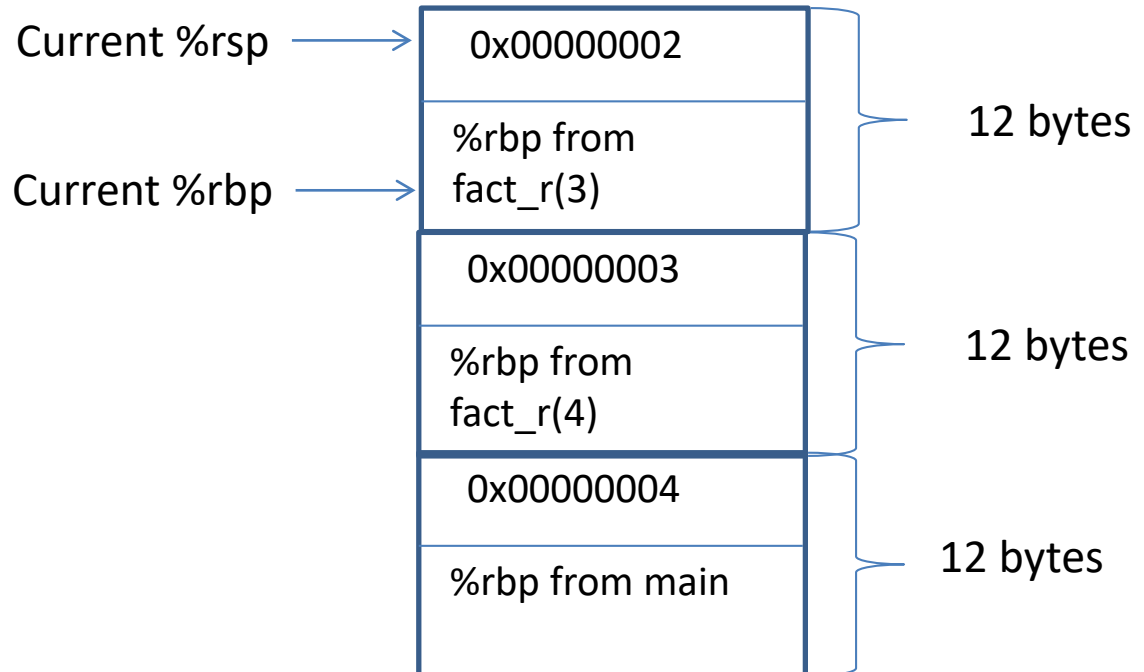
.L2:

```
movl  -4(%rbp), %eax
subl  $1, %eax
movl  %eax, %edi
call  fact_r
imull -4(%rbp), %eax
```

.L3:

```
leave
ret
```

fact_r(2):



fact_r:

```
pushq %rbp
movq  %rsp, %rbp
subq  $16, %rsp
movl  %edi, -4(%rbp)
cmpl  $1, -4(%rbp)
jg    .L2
movl  $1, %eax
jmp   .L3
```

.L2:

```
movl  -4(%rbp), %eax
subl  $1, %eax
movl  %eax, %edi
call  fact_r
imull -4(%rbp), %eax
```


.L3:

```
leave
ret
```

fact_r(1):

Current %rsp →

Current %rbp →



```
int fact_r(int x) {  
    if (x <= 1)  
        return 1;  
    else return x * fact_r(x-1);  
}
```

Where is x remembered?

Now, using

```
gcc -o fact.s -s fact.c -O2
```

fact_r:

 cmpl \$1, %edi

 movl \$1, %eax

 jg .L3

 jmp .L2

.L9:

 movl %edx, %edi

.L3:

 leal -1(%rdi), %edx

 imull %edi, %eax

 cmpl \$1, %edx

 jne .L9

.L2:

 ret

**This is no longer
recursive!!**

**Compiler removes
recursion whenever
possible, because loops
are faster**

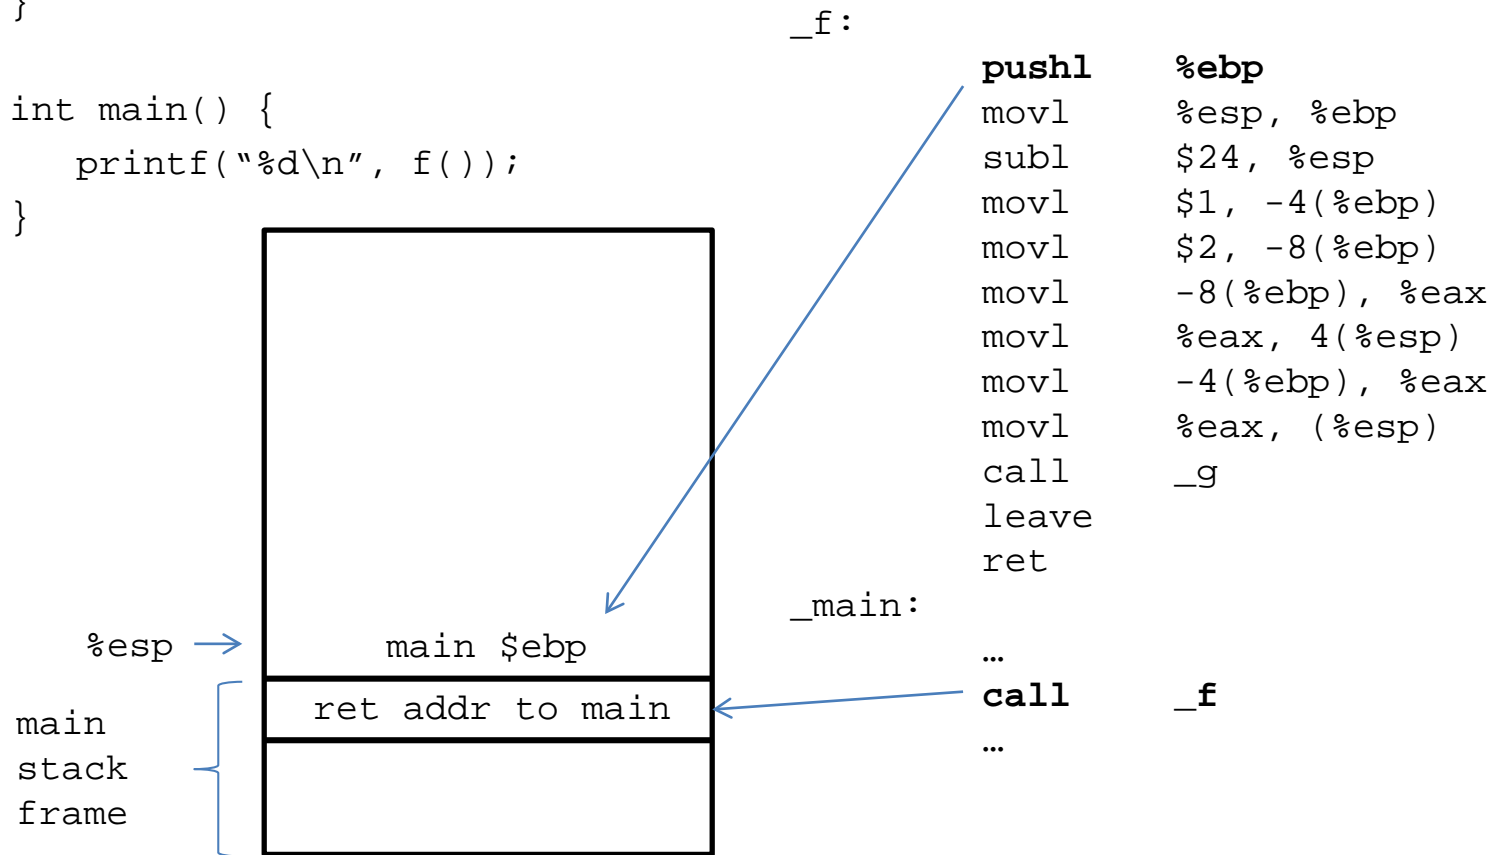
Function parameters and the program stack

- In x86-64, parameters are usually passed through registers
- In the unlikely event that more than 6 parameters are passed, then the program stack is used for the rest of the parameters
- To illustrate, we will switch to IA-32, and consider this code.

```
int g(int x, int y) {  
    return x + y;  
}  
  
int f() {  
    int a=1, b=2;  
    return g(a,b);  
}  
  
int main() {  
    printf("%d\n", f());  
}
```

```
int f() {
    int a=1, b=2;
    return g(a,b);
}
```

```
int main() {
    printf("%d\n", f());
}
```

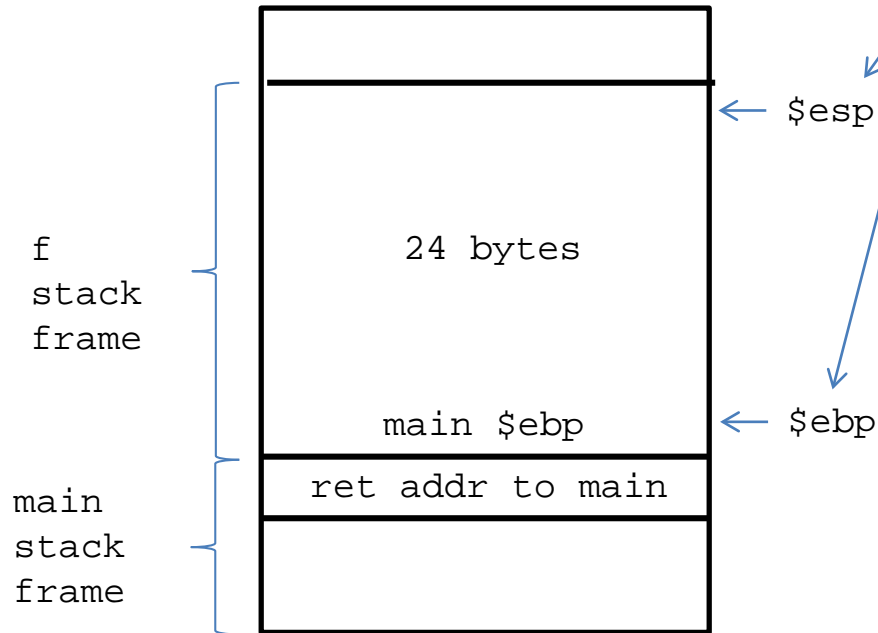



```
int f() {
    int a=1, b=2;
    return g(a,b);
}
```

```
int main() {
    printf("%d\n", f());
}
```

_f:

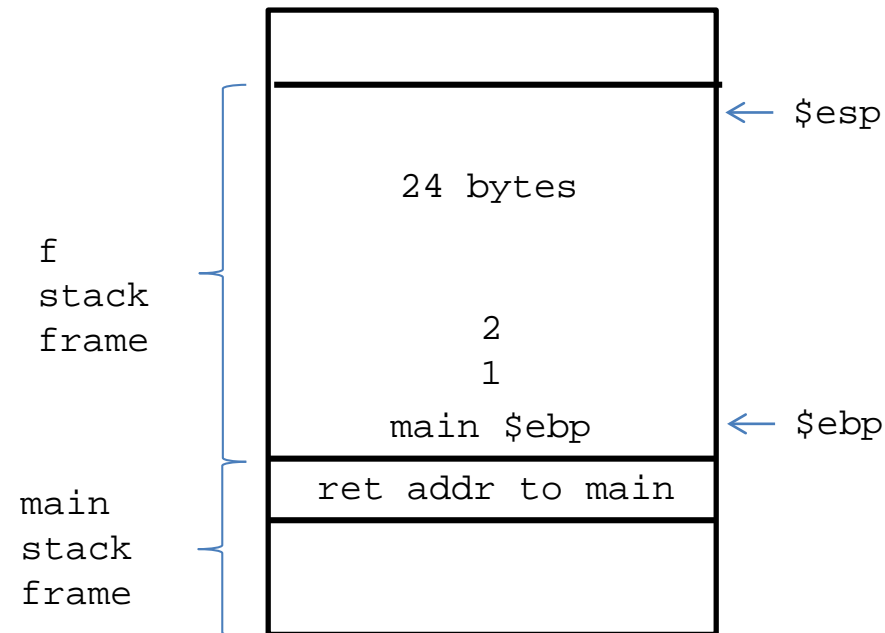
```
pushl    %ebp
movl     %esp, %ebp
subl     $24, %esp
movl     $1, -4(%ebp)
movl     $2, -8(%ebp)
movl     -8(%ebp), %eax
movl     %eax, 4(%esp)
movl     -4(%ebp), %eax
movl     %eax, (%esp)
call     _g
leave
ret
```



```
int f() {
    int a=1, b=2;
    return g(a,b);
}
```

```
int main() {
    printf("%d\n", f());
}
```

_f:



```
pushl    %ebp
movl     %esp, %ebp
subl     $24, %esp
movl     $1, -4(%ebp)
movl     $2, -8(%ebp)
movl     -8(%ebp), %eax
movl     %eax, 4(%esp)
movl     -4(%ebp), %eax
movl     %eax, (%esp)
call     _g
leave
ret
```

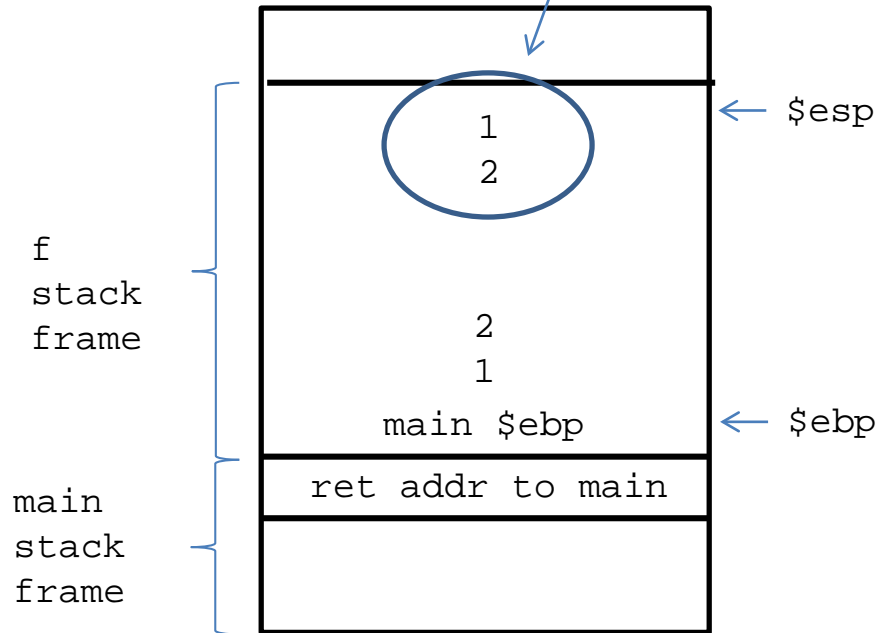
```
int f() {
    int a=1, b=2;
    return g(a,b);
}
```

These are parameters that will be passed to g

```
int main() {
    printf("%d\n", f());
}
```

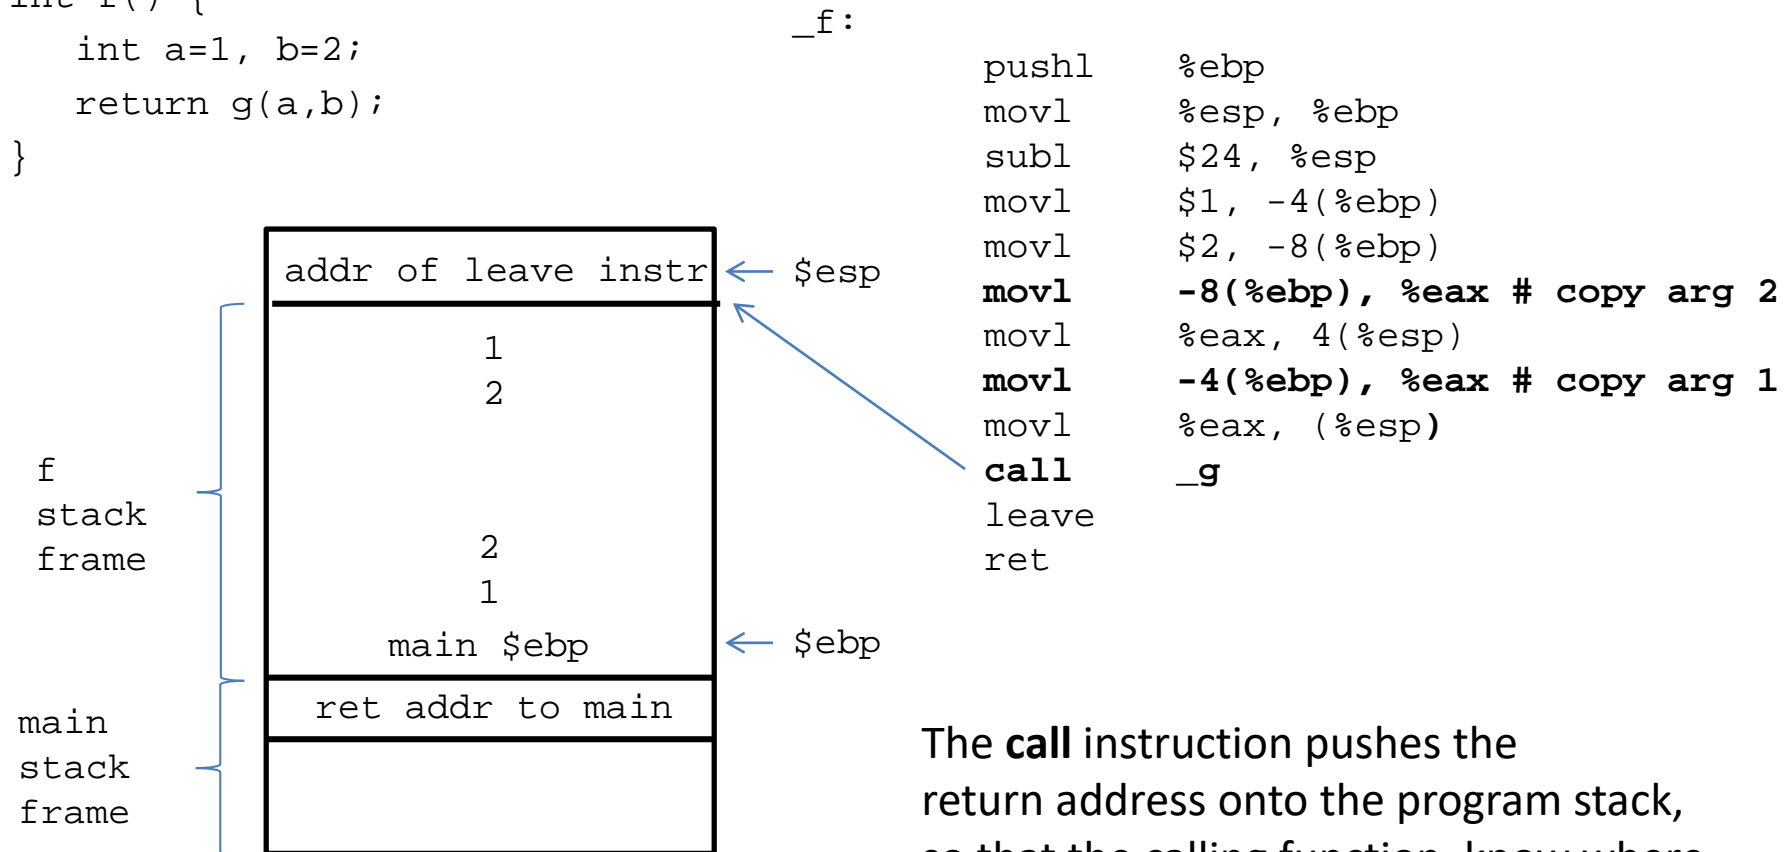
`_f:`

```
pushl    %ebp
movl     %esp, %ebp
subl     $24, %esp
movl     $1, -4(%ebp)
movl     $2, -8(%ebp)
movl     -8(%ebp), %eax
movl     %eax, 4(%esp)
movl     -4(%ebp), %eax
movl     %eax, (%esp)
call     _g
leave
ret
```



```
int g(int x, int y) {
    return x + y;
}
```

```
int f() {
    int a=1, b=2;
    return g(a,b);
}
```



The **call** instruction pushes the return address onto the program stack, so that the calling function know where to return

```
int g(int x, int y) {
    return x + y;
}
```

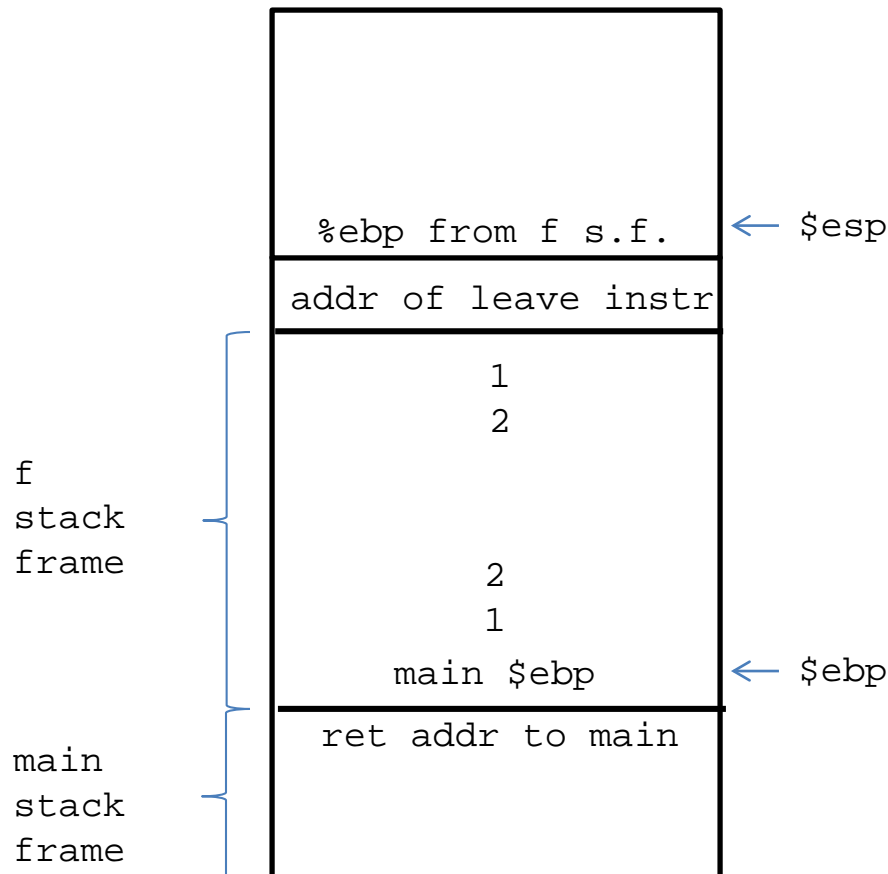
```
int f() {
    int a=1, b=2;
    return g(a,b);
}
```

`_g:`

```
pushl    %ebp
movl     %esp, %ebp
movl     12(%ebp), %eax
addl     8(%ebp), %eax
popl     %ebp
ret
```

`_f:`

```
pushl    %ebp
movl     %esp, %ebp
subl     $24, %esp
movl     $1, -4(%ebp)
movl     $2, -8(%ebp)
movl     -8(%ebp), %eax
movl     %eax, 4(%esp)
movl     -4(%ebp), %eax
movl     %eax, (%esp)
call     _g
leave
ret
```



```
int g(int x, int y) {
    return x + y;
}
```


```
int f() {
    int a=1, b=2;
    return g(a,b);
}
```

`_g:`

```
pushl    %ebp
movl    %esp, %ebp
movl     12(%ebp), %eax
addl     8(%ebp), %eax
popl     %ebp
ret
```

`_f:`

```
pushl    %ebp
movl     %esp, %ebp
subl     $24, %esp
movl     $1, -4(%ebp)
movl     $2, -8(%ebp)
movl     -8(%ebp), %eax
movl     %eax, 4(%esp)
movl     -4(%ebp), %eax
movl     %eax, (%esp)
call     _g
leave
ret
```



```
int g(int x, int y) {
    return x + y;
}
```

`_g:`

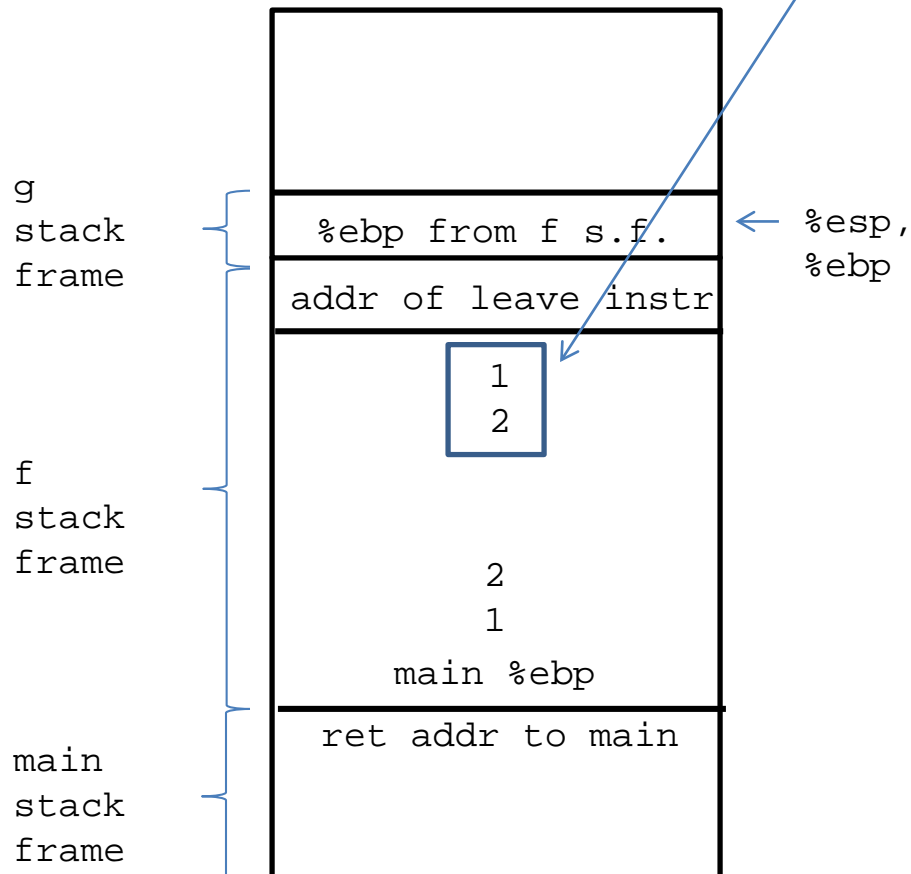
```
pushl    %ebp
movl     %esp, %ebp
movl     -12(%ebp), %eax
addl     -8(%ebp), %eax
popl     %ebp
ret
```

```
int f() {
    int a=1, b=2;
    return g(a,b);
}
```

`-12(%ebp)` and `-8(%ebp)`
Are the parameters
Passed to g

`_f:`

```
pushl    %ebp
movl     %esp, %ebp
subl     $24, %esp
movl     $1, -4(%ebp)
movl     $2, -8(%ebp)
movl     -8(%ebp), %eax
movl     %eax, 4(%esp)
movl     -4(%ebp), %eax
movl     %eax, (%esp)
call     _g
leave
ret
```



```
int g(int x, int y) {                                _g:
    return x + y;
}
```

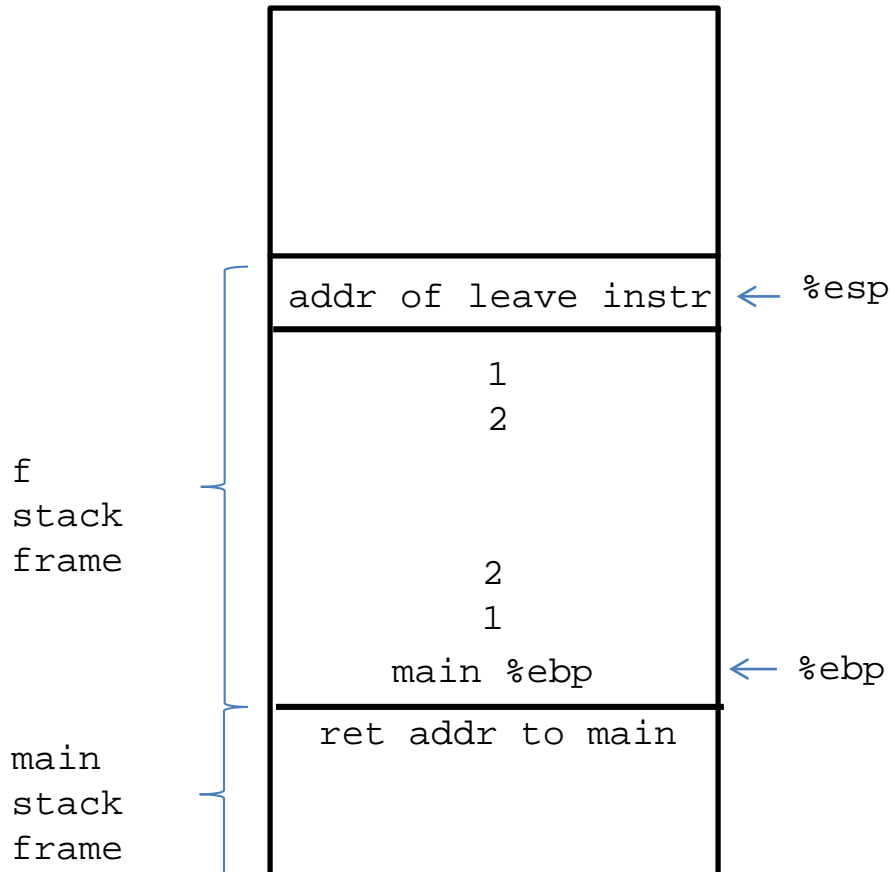
```
int f() {
    int a=1, b=2;
    return g(a,b);
}
```

12(%ebp) and 8(%ebp)
Are the parameters
Passed to g

_f:

```
pushl    %ebp
movl     %esp, %ebp
movl     12(%ebp), %eax
addl     8(%ebp), %eax
popl    %ebp
ret
```

```
pushl    %ebp
movl     %esp, %ebp
subl     $24, %esp
movl     $1, -4(%ebp)
movl     $2, -8(%ebp)
movl     -8(%ebp), %eax
movl     %eax, 4(%esp)
movl     -4(%ebp), %eax
movl     %eax, (%esp)
call     _g
leave
ret
```




```
int g(int x, int y) {
    return x + y;
}
```

`_g:`

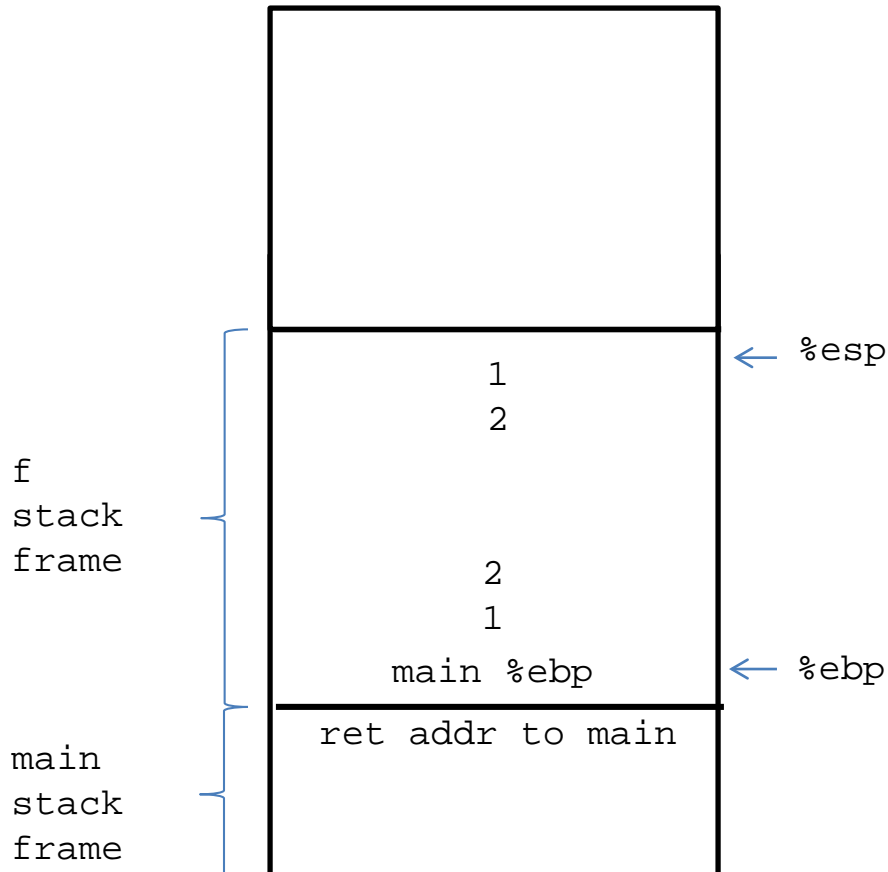
```
pushl    %ebp
movl     %esp, %ebp
movl     12(%ebp), %eax
addl     8(%ebp), %eax
popl     %ebp
ret
```

```
int f() {
    int a=1, b=2;
    return g(a,b);
}
```

Next instruction is
"leave" in f

`_f:`

```
pushl    %ebp
movl     %esp, %ebp
subl     $24, %esp
movl     $1, -4(%ebp)
movl     $2, -8(%ebp)
movl     -8(%ebp), %eax
movl     %eax, 4(%esp)
movl     -4(%ebp), %eax
movl     %eax, (%esp)
call     _g
leave
ret
```




```
int g(int x, int y) {
    return x + y;
}
```

```
int f() {
    int a=1, b=2;
    return g(a,b);
}
```

`_f:`

```
pushl    %ebp
movl     %esp, %ebp
subl     $24, %esp
movl     $1, -4(%ebp)
movl     $2, -8(%ebp)
movl     -8(%ebp), %eax
movl     %eax, 4(%esp)
movl     -4(%ebp), %eax
movl     %eax, (%esp)
call     _g
leave
ret
```



leave is a macro for popping into `%ebp`. After **ret**, the next instruction is back in **main**


```
int g(int x, int y) {
    return x + y;
}
```

```
int f() {
    int a=1, b=2;
    return g(a,b);
}
```

`_f:`

```
pushl    %ebp
movl     %esp, %ebp
subl     $24, %esp
movl     $1, -4(%ebp)
movl     $2, -8(%ebp)
movl     -8(%ebp), %eax
movl     %eax, 4(%esp)
movl     -4(%ebp), %eax
movl     %eax, (%esp)
call     _g
leave
ret
```

main
stack
frame



← \$esp

← \$ebp

Example illustrating program stack

```
int main( ) {
    int x, y, z;
    printf("Type 3 numbers\n");
    scanf("%d%d%d", &x, &y, &z);
    printf("%d+%d+%d=%d\n", x, y,
           z, sum3(x, y, z));
}

int sum3(int x, int y, int z) {
    return x + sum2(y,z);
}

int sum2(int x, int y) {
    int sum = x + y;
    return sum;
}
```

x86-64 Example

```
int main( ) {  
    int x, y, z;  
    printf("Type 3 numbers\n");  
    scanf("%d%d%d", &x, &y, &z);  
    printf("%d+%d+%d=%d\n", x, y,  
          z, sum3(x, y, z));  
}
```

main:

```
    pushq    %rbp  
    movq     %rsp, %rbp  
    subq     $16, %rsp  
    movl     $.LC0, %edi  
    call     puts  
    leaq     -16(%rbp), %rcx  
    leaq     -12(%rbp), %rdx  
    leaq     -8(%rbp), %rsi
```


```
    movl     $.LC1, %edi  
    call     __isoc99_scanf  
    movl     -16(%rbp), %edx  
    movl     -12(%rbp), %ecx  
    movl     -8(%rbp), %eax  
    movl     %ecx, %esi  
    movl     %eax, %edi  
    call     sum3  
    movl     %eax, -4(%rbp)  
    movl     -4(%rbp), %eax  
    movl     %eax, %esi  
    movl     $.LC2, %edi  
    call     printf  
    leave  
    ret
```

x86-64 Example

main:

```
pushq    %rbp
movq     %rsp, %rbp
subq     $16, %rsp
movl     $.LC0, %edi
call     puts
leaq     -16(%rbp), %rcx
leaq     -12(%rbp), %rdx
leaq     -8(%rbp), %rsi
movl     $.LC1, %edi
call     __isoc99_scanf
movl     -16(%rbp), %edx
movl     -12(%rbp), %ecx
movl     -8(%rbp), %eax
movl     %ecx, %esi
movl     %eax, %edi
call     sum3
```

```
movl     %eax, -4(%rbp)
movl     -4(%rbp), %eax
movl     %eax, %esi
movl     $.LC2, %edi
call     printf
leave
ret
```




24 byte
Stack frame
For main

x86-64 Example

main:

```
pushq    %rbp
movq     %rsp, %rbp
subq     $16, %rsp
movl     $.LC0, %edi
call     puts
leaq     -16(%rbp), %rcx
leaq     -12(%rbp), %rdx
leaq     -8(%rbp), %rsi
movl     $.LC1, %edi
call     __isoc99_scanf
movl     -16(%rbp), %edx
movl     -12(%rbp), %ecx
movl     -8(%rbp), %eax
movl     %ecx, %esi
movl     %eax, %edi
call     sum3
```

```
movl     %eax, -4(%rbp)
movl     -4(%rbp), %eax
movl     %eax, %esi
movl     $.LC2, %edi
call     printf
leave
ret
```

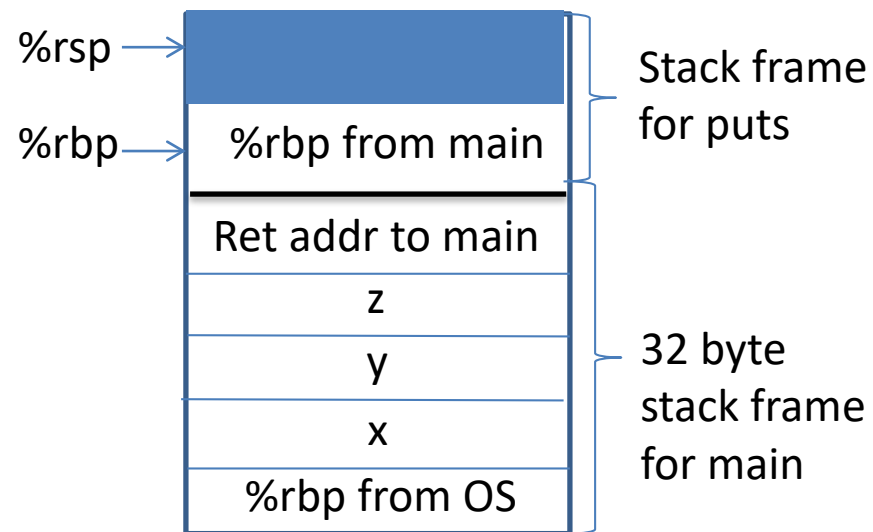


x86-64 Example

main:

```
    pushq    %rbp
    movq     %rsp, %rbp
    subq     $16, %rsp
Arg to { movl    $.LC0, %edi
puts { call    puts
    leaq     -16(%rbp), %rcx
    leaq     -12(%rbp), %rdx
    leaq     -8(%rbp), %rsi
    movl     $.LC1, %edi
    call     __isoc99_scanf
    movl     -16(%rbp), %edx
    movl     -12(%rbp), %ecx
    movl     -8(%rbp), %eax
    movl     %ecx, %esi
    movl     %eax, %edi
    call     sum3
```

```
    movl     %eax, -4(%rbp)
    movl     -4(%rbp), %eax
    movl     %eax, %esi
    movl     $.LC2, %edi
    call     printf
    leave
    ret
```




x86-64 Example

main:

```
pushq    %rbp
movq     %rsp, %rbp
subq     $16, %rsp
movl     $.LC0, %edi
call     puts    # prompt user
leaq     -16(%rbp), %rcx # &z
leaq     -12(%rbp), %rdx # &y
leaq     -8(%rbp), %rsi  # &x
movl     $.LC1, %edi
call     __isoc99_scanf
movl     -16(%rbp), %edx
movl     -12(%rbp), %ecx
movl     -8(%rbp), %eax
movl     %ecx, %esi
movl     %eax, %edi
call     sum3
```

```
movl     %eax, -4(%rbp)
movl     -4(%rbp), %eax
movl     %eax, %esi
movl     $.LC2, %edi
call     printf
leave
ret
```



x86-64 Example

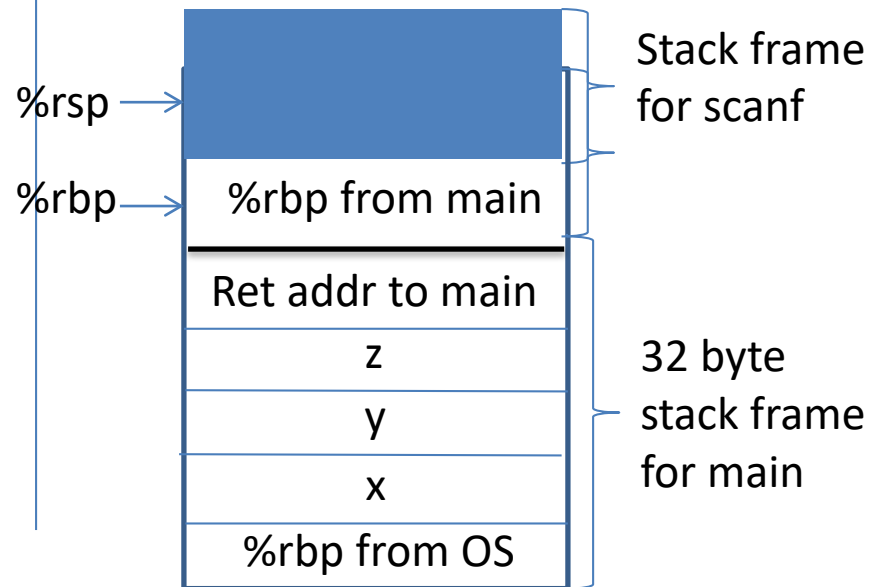
main:

```
pushq    %rbp
movq     %rsp, %rbp
subq     $16, %rsp
movl     $.LC0, %edi
call     puts    # prompt user
leaq     -16(%rbp), %rcx
leaq     -12(%rbp), %rdx
leaq     -8(%rbp), %rsi
movl     $.LC1, %edi
call     __isoc99_scanf
movl     -16(%rbp), %edx
movl     -12(%rbp), %ecx
movl     -8(%rbp), %eax
movl     %ecx, %esi
movl     %eax, %edi
call     sum3
```

args to
scanf



```
movl     %eax, -4(%rbp)
movl     -4(%rbp), %eax
movl     %eax, %esi
movl     $.LC2, %edi
call     printf
leave
ret
```

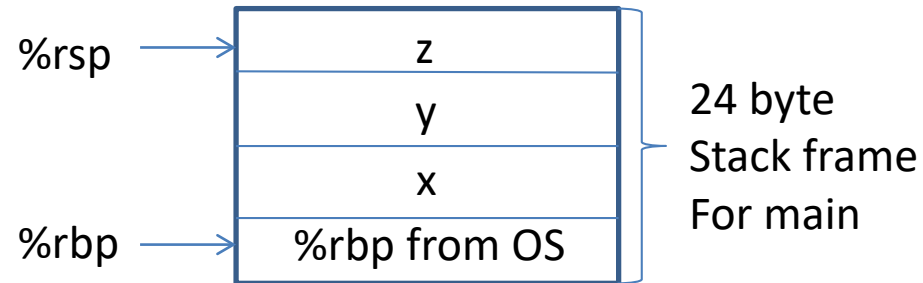


x86-64 Example

main:

```
pushq    %rbp
movq     %rsp, %rbp
subq     $16, %rsp    #%rbp pushed
movl     $.LC0, %edi
call     puts    # prompt user
leaq     -16(%rbp), %rcx
leaq     -12(%rbp), %rdx
leaq     -8(%rbp), %rsi
movl     $.LC1, %edi
call     __isoc99_scanf
movl     -16(%rbp), %edx #z
movl     -12(%rbp), %ecx
movl     -8(%rbp), %eax
movl     %ecx, %esi    #y
movl     %eax, %edi    #z
call     sum3
```

```
movl     %eax, -4(%rbp)
movl     -4(%rbp), %eax
movl     %eax, %esi
movl     $.LC2, %edi
call     printf
leave
ret
```



x86-64 Example

```
movl    -16(%rbp), %edx #z
movl    -12(%rbp), %ecx
movl    -8(%rbp), %eax
movl    %ecx, %esi      #y
movl    %eax, %edi      #z
call    sum3
```

(gdb) x/x \$rsp

0x7fffffff978: 0x00400634

In main:

0x0040062f <+66>: callq 0x40064d <sum3>

0x00400634 <+71>: mov %eax,-0x4(%rbp)


When one function calls another,
The return address in the calling
function is pushed onto the program
stack

(gdb) break *sum3

Breakpoint 5 at 0x40064d

(gdb) cont

Continuing.

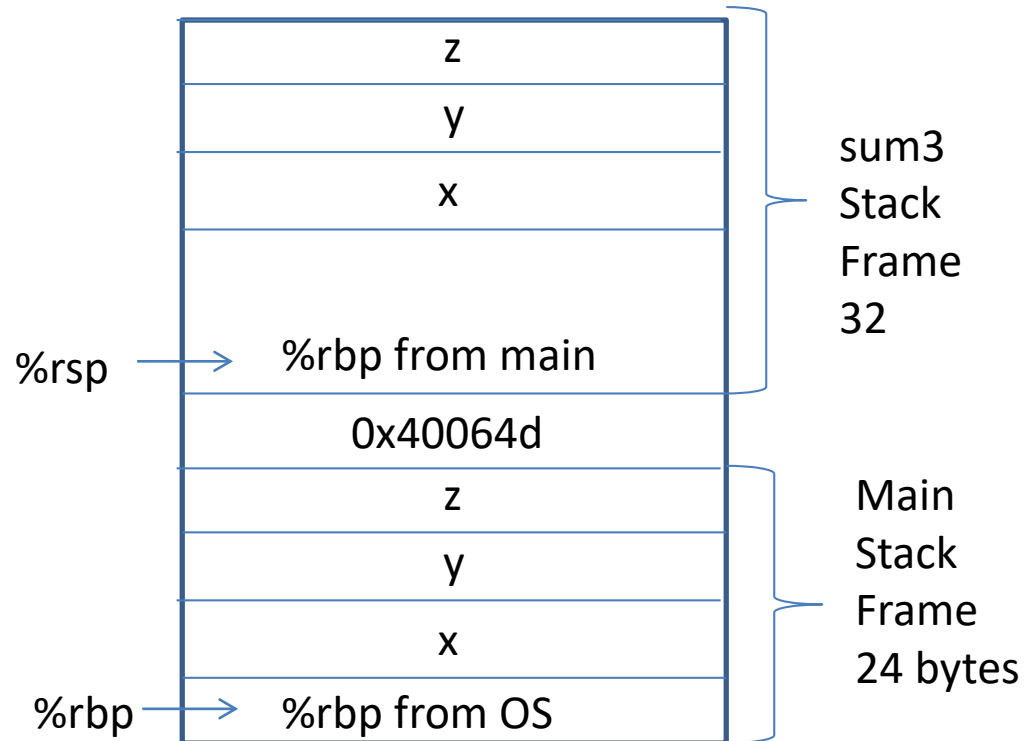


Breakpoint 5, 0x00000000004006

x86-64 Example

sum3:

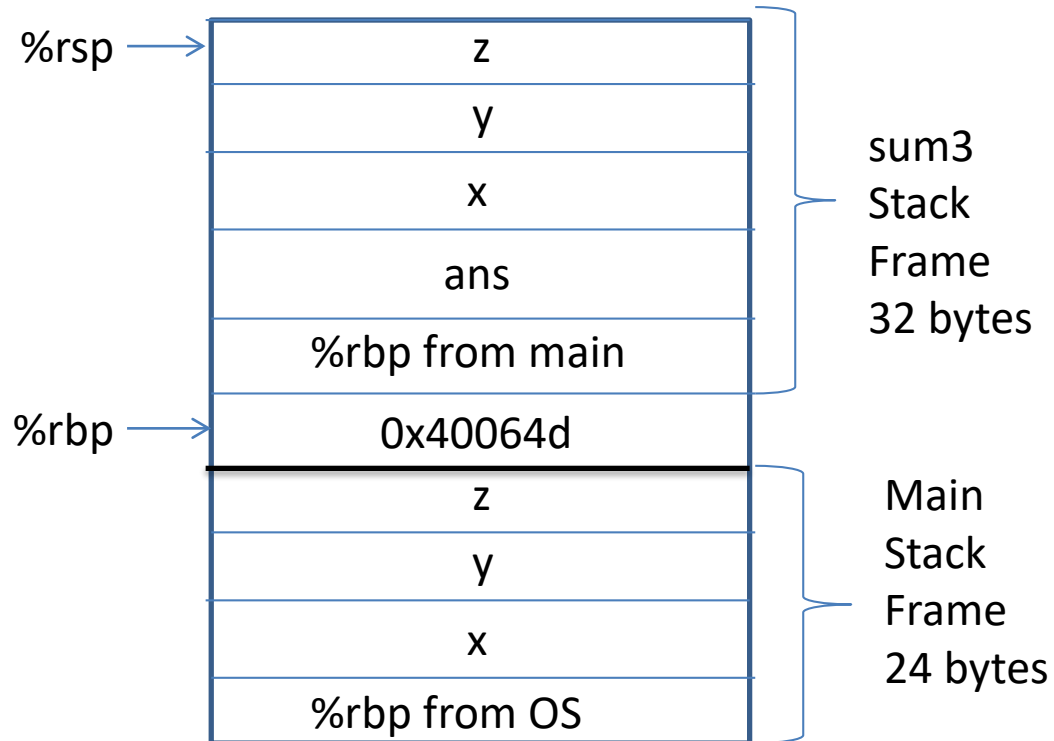
```
pushq    %rbp
movq     %rsp, %rbp
.cfi_def_cfa_register 6
subq     $32, %rsp
movl     %edi, -20(%rbp)
movl     %esi, -24(%rbp)
movl     %edx, -28(%rbp)
movl     -28(%rbp), %edx
movl     -24(%rbp), %eax
movl     %edx, %esi
movl     %eax, %edi
call     sum2
movl     %eax, -4(%rbp)
movl     -4(%rbp), %eax
movl     -20(%rbp), %edx
addl     %edx, %eax
leave
ret
```



x86-64 Example

sum3:

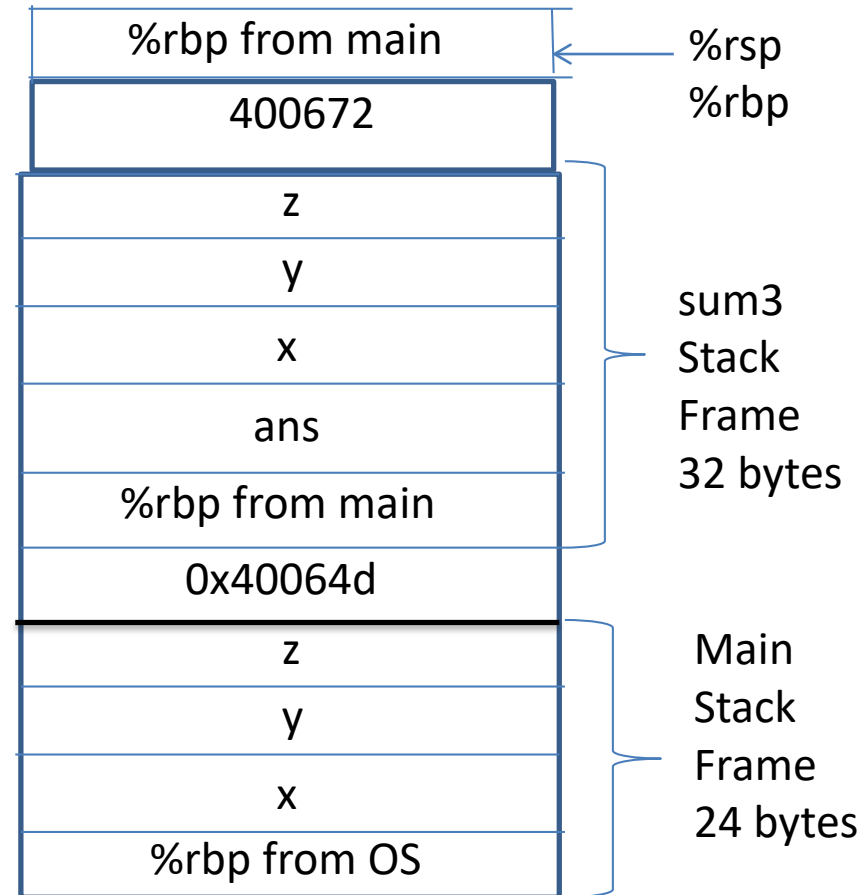
```
pushq    %rbp
movq     %rsp, %rbp
subq    $32, %rsp
movl    %edi, -20(%rbp)
movl    %esi, -24(%rbp)
movl    %edx, -28(%rbp)
movl     -28(%rbp), %edx
movl     -24(%rbp), %eax
movl     %edx, %esi
movl     %eax, %edi
call    sum2
movl     %eax, -4(%rbp)
movl     -4(%rbp), %eax
movl     -20(%rbp), %edx
addl     %edx, %eax
leave
ret
```



x86-64 Example

sum2:

```
0x40067f <+0>:  push    %rbp
0x400680 <+1>:  mov     %rsp,%rbp
0x400683 <+4>:  mov     %edi,-0x4(%rbp)
0x400686 <+7>:  mov     %esi,-0x8(%rbp)
0x400689 <+10>:  mov     -0x8(%rbp),%eax
0x40068c <+13>:  mov     -0x4(%rbp),%edx
0x40068f <+16>:  add     %edx,%eax
0x400691 <+18>:  pop     %rbp
0x400692 <+19>:  retq
```



The GNU debugger

- Very low-level debugging / code inspection tool
- Operates on executable files
- Enables programmers to “disassemble” code, set breakpoints, inspect register and memory contents, step through code execution, etc.
- We will use it for code inspection, and to help us understand what machine code is doing and how
- We will also use it to understand a certain type of security vulnerability caused by the stack frame organization

Example gdb session

```
> gdb sum3
```

```
> disas sum3
```

```
(gdb) disas sum3
```

```
Dump of assembler code for function sum3:
```

```
0x0000000000400608 <+0>:      push    %rbp
0x0000000000400609 <+1>:      mov     %rsp,%rbp
0x000000000040060c <+4>:      sub     $0x20,%rsp
0x0000000000400610 <+8>:      mov     %rdi,-0x18(%rbp)
0x0000000000400614 <+12>:     mov     -0x18(%rbp),%rax
0x0000000000400618 <+16>:     mov     (%rax),%eax
0x000000000040061a <+18>:     mov     %eax,-0x4(%rbp)
0x000000000040061d <+21>:     mov     -0x18(%rbp),%rax
0x0000000000400621 <+25>:     add     $0x4,%rax
0x0000000000400625 <+29>:     mov     %rax,%rdi
0x0000000000400628 <+32>:     mov     $0x0,%eax
0x000000000040062d <+37>:     callq   0x40063a <sum2>
0x0000000000400632 <+42>:     add     %eax,-0x4(%rbp)
0x0000000000400635 <+45>:     mov     -0x4(%rbp),%eax
0x0000000000400638 <+48>:     leaveq
0x0000000000400639 <+49>:     retq
```

Example gdb session

```
(gdb) run
```

```
Starting program: /home/DPU/slytinen/406w18/x86/sum3_slow
```

```
Type 3 integers
```

```
1 2 3
```

```
Breakpoint 1, 0x0000000000400608 in sum3 ()
```

```
Missing separate debuginfos, use: debuginfo-install glibc-2.17-106.el7_2.6.x86_64
```

```
(gdb) print/x $rdi
```

```
$1 = 0x7fffffff960
```

```
(gdb) x/3x $rdi
```

```
0x7fffffff960: 0x00000001      0x00000002      0x00000003
```

```
(gdb) print/x $rsp
```

```
$2 = 0x7fffffff958
```

```
(gdb) x/x $rsp
```

```
0x7fffffff958: 0x004005ef
```

```
(gdb) x/x 0x4005ef
```

```
0x4005ef <main+79>:      0x8bf84589
```

```
(gdb) x/i main+79
```

```
0x4005ef <main+79>:  mov    %eax,-0x8(%rbp)
```

```
(gdb) x/i main+74
```

```
0x4005ea <main+74>:  callq 0x400608 <sum3>
```

Example gdb session

```
(gdb) break *sum3+48
```

```
Breakpoint 2 at 0x400638
```

```
(gdb) continue
```

```
Continuing.
```

```
Breakpoint 2, 0x0000000000400638 in sum3 ()
```

```
(gdb) print/d $rax
```

```
$3 = 6
```

```
(gdb) stepi
```

```
0x0000000000400639 in sum3 ()
```

```
(gdb) stepi
```

```
0x00000000004005ef in main ()
```

```
(gdb) disas 0x4005f2
```

```
0x00000000004005ea <+74>:    callq   0x400608 <sum3>
0x00000000004005ef <+79>:    mov     %eax,-0x8(%rbp)
=> 0x00000000004005f2 <+82>:    mov     -0x8(%rbp),%eax
0x00000000004005f5 <+85>:    mov     %eax,%esi
0x00000000004005f7 <+87>:    mov     $0x400713,%edi
0x00000000004005fc <+92>:    mov     $0x0,%eax
0x0000000000400601 <+97>:    callq   0x400470 <printf@plt>
0x0000000000400606 <+102>:   leaveq
0x0000000000400607 <+103>:   retq
```

Example gdb session

```
(gdb) break *main+97
```

```
Breakpoint 3 at 0x400601
```

```
(gdb) continue
```

```
Continuing.
```

```
Breakpoint 3, 0x0000000000400601 in main ()
```

```
(gdb) print/x $rdi
```

```
$4 = 0x400713
```

```
(gdb) x/x $rdi
```

```
0x400713:          0x20656854
```

```
(gdb) x/s $rdi
```

```
0x400713:          "The sum is %d\n"
```

```
(gdb) print/d $rsi
```

```
$5 = 6
```

```
(gdb) cont
```

```
Continuing.
```

```
The sum is 6
```

```
[Inferior 1 (process 27473) exited with code 015]
```

Basic gdb commands

- Start the debugger: **`gdb`** followed by the name of the executable
- **`disas`**: “disassemble” a function (translate from machine code back to assembly language). Can either be followed by a function name or an address
- **`break`**: set a breakpoint. Can either be followed by an address or an offset from the beginning of a function. Place an `*` beforehand, such as `break *main`
- **`run`**
- **`cont`**: resume execution from a breakpoint
- **`print`**: print the contents of a register. Qualifiers: `/d`, `/x`, `/s`, `/c`
 - Example: **`print/d $rax`**
- **`x`**: print the contents of a memory address. Also may use qualifiers
 - Example: `x/x 0x400601`
 - `x/3x $rdi`
 - `x/s $rdi`
- **`stepi`**: execute the next instruction (step into). If the current instruction is a function call, then the first instruction of the called function is executed